

Disclaimer: It is probably easier to read this assignment as .Rmd and not as PDF, because there are a lot of outputs and plots, which can get a little overwhelming when not viewed through RStudio.

Exercise 1

In this exercise we use the College data set from the ISLR package. We predict the number of applications received using the other variables.

```
library(ISLR)
library(tidymodels)

library(tidymodels)
library(funModeling)
library(ISLR)
library(vip)
library(forcats)
library(GGally)

?ISLR::College

College <- tibble(College)
College

## # A tibble: 777 x 18
##   Private Apps Accept Enroll Top10perc Top25perc F.Undergrad P.Undergrad
##   <fct>   <dbl> <dbl> <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 Yes     1660  1232   721      23      52     2885     537
## 2 Yes     2186  1924   512      16      29     2683    1227
## 3 Yes     1428  1097   336      22      50     1036      99
## 4 Yes      417   349   137      60      89      510      63
## 5 Yes      193   146    55      16      44      249     869
## 6 Yes      587   479   158      38      62      678      41
## 7 Yes      353   340   103      17      45      416     230
## 8 Yes     1899  1720   489      37      68     1594      32
## 9 Yes     1038   839   227      30      63      973     306
## 10 Yes      582   498   172      21      44      799      78
## # ... with 767 more rows, and 10 more variables: Outstate <dbl>,
## #   Room.Board <dbl>, Books <dbl>, Personal <dbl>, PhD <dbl>, Terminal <dbl>,
## #   S.F.Ratio <dbl>, perc.alumni <dbl>, Expend <dbl>, Grad.Rate <dbl>

basic_eda <- function(data) {
  glimpse(data)
  print(status(data))
  freq(data)
  print(profiling_num(data))
  plot_num(data)
  describe(data)
}

basic_eda(College)

## Rows: 777
## Columns: 18
## $ Private      <fct> Yes, Yes, Yes, Yes, Yes, Yes, Yes, Yes, Yes, Yes, Yes, Yes~
## $ Apps         <dbl> 1660, 2186, 1428, 417, 193, 587, 353, 1899, 1038, 582, 173~
```

```

## $ Accept      <dbl> 1232, 1924, 1097, 349, 146, 479, 340, 1720, 839, 498, 1425~
## $ Enroll      <dbl> 721, 512, 336, 137, 55, 158, 103, 489, 227, 172, 472, 484,~
## $ Top10perc   <dbl> 23, 16, 22, 60, 16, 38, 17, 37, 30, 21, 37, 44, 38, 44, 23~
## $ Top25perc   <dbl> 52, 29, 50, 89, 44, 62, 45, 68, 63, 44, 75, 77, 64, 73, 46~
## $ F.Undergrad <dbl> 2885, 2683, 1036, 510, 249, 678, 416, 1594, 973, 799, 1830~
## $ P.Undergrad <dbl> 537, 1227, 99, 63, 869, 41, 230, 32, 306, 78, 110, 44, 638~
## $ Outstate    <dbl> 7440, 12280, 11250, 12960, 7560, 13500, 13290, 13868, 1559~
## $ Room.Board  <dbl> 3300, 6450, 3750, 5450, 4120, 3335, 5720, 4826, 4400, 3380~
## $ Books       <dbl> 450, 750, 400, 450, 800, 500, 500, 450, 300, 660, 500, 400~
## $ Personal    <dbl> 2200, 1500, 1165, 875, 1500, 675, 1500, 850, 500, 1800, 60~
## $ PhD         <dbl> 70, 29, 53, 92, 76, 67, 90, 89, 79, 40, 82, 73, 60, 79, 36~
## $ Terminal    <dbl> 78, 30, 66, 97, 72, 73, 93, 100, 84, 41, 88, 91, 84, 87, 6~
## $ S.F.Ratio   <dbl> 18.1, 12.2, 12.9, 7.7, 11.9, 9.4, 11.5, 13.7, 11.3, 11.5, ~
## $ perc.alumni <dbl> 12, 16, 30, 37, 2, 11, 26, 37, 23, 15, 31, 41, 21, 32, 26,~
## $ Expend      <dbl> 7041, 10527, 8735, 19016, 10922, 9727, 8861, 11487, 11644,~
## $ Grad.Rate   <dbl> 60, 56, 54, 59, 15, 55, 63, 73, 80, 52, 73, 76, 74, 68, 55~
##               variable q_zeros  p_zeros q_na p_na q_inf p_inf  type
## Private      Private    0 0.000000000    0  0    0    0  factor
## Apps         Apps      0 0.000000000    0  0    0    0  numeric
## Accept       Accept    0 0.000000000    0  0    0    0  numeric
## Enroll       Enroll    0 0.000000000    0  0    0    0  numeric
## Top10perc    Top10perc  0 0.000000000    0  0    0    0  numeric
## Top25perc    Top25perc  0 0.000000000    0  0    0    0  numeric
## F.Undergrad  F.Undergrad 0 0.000000000    0  0    0    0  numeric
## P.Undergrad  P.Undergrad 0 0.000000000    0  0    0    0  numeric
## Outstate     Outstate   0 0.000000000    0  0    0    0  numeric
## Room.Board   Room.Board  0 0.000000000    0  0    0    0  numeric
## Books        Books      0 0.000000000    0  0    0    0  numeric
## Personal     Personal    0 0.000000000    0  0    0    0  numeric
## PhD          PhD        0 0.000000000    0  0    0    0  numeric
## Terminal     Terminal    0 0.000000000    0  0    0    0  numeric
## S.F.Ratio    S.F.Ratio   0 0.000000000    0  0    0    0  numeric
## perc.alumni  perc.alumni 2 0.002574003    0  0    0    0  numeric
## Expend       Expend     0 0.000000000    0  0    0    0  numeric
## Grad.Rate    Grad.Rate   0 0.000000000    0  0    0    0  numeric
##               unique
## Private      2
## Apps         711
## Accept       693
## Enroll       581
## Top10perc    82
## Top25perc    89
## F.Undergrad  714
## P.Undergrad  566
## Outstate     640
## Room.Board   553
## Books        122
## Personal     294
## PhD          78
## Terminal     65
## S.F.Ratio    173
## perc.alumni  61
## Expend       744
## Grad.Rate    81

```

```
## Warning: `guides(<scale> = FALSE)` is deprecated. Please use `guides(<scale> =`  
## "none")` instead.
```

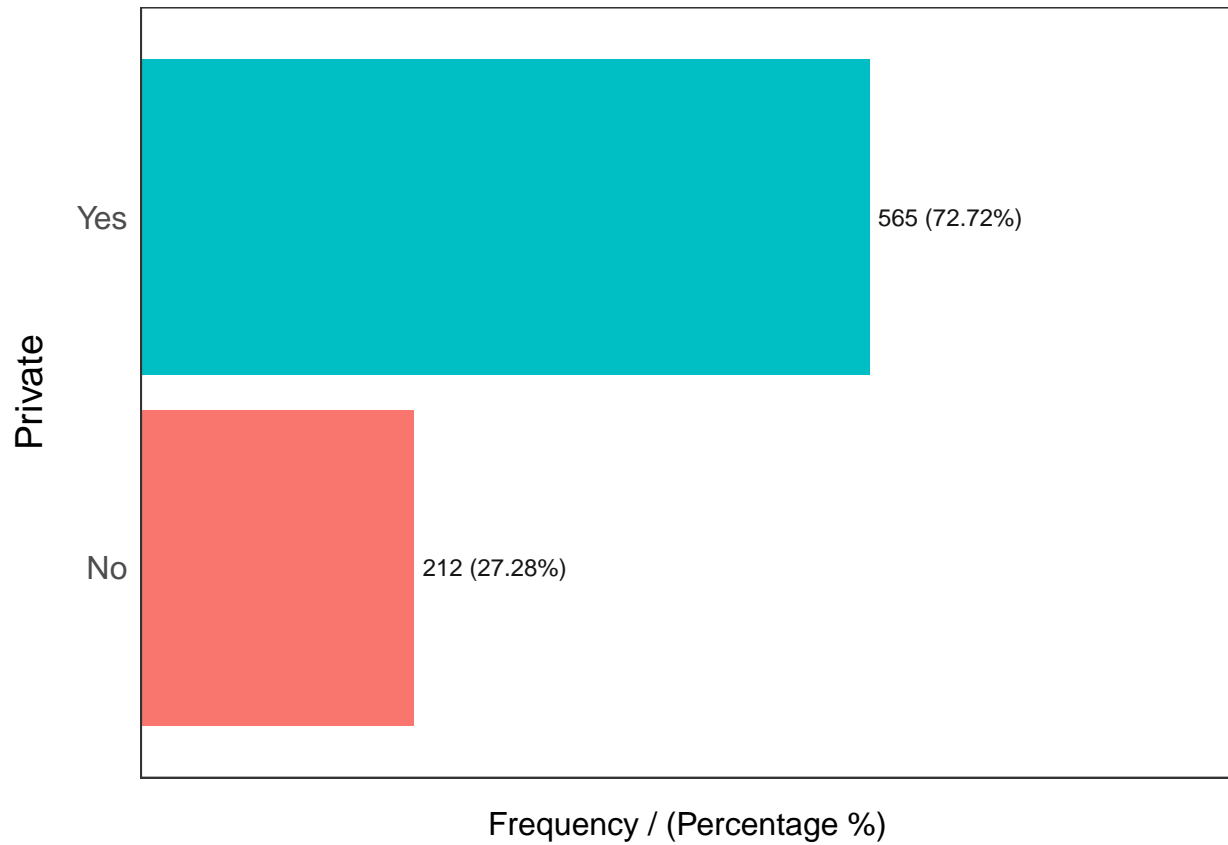
##	variable	mean	std_dev	variation_coef	p_01	p_05	p_25
## 1	Apps	3001.63835	3870.201484	1.2893630	192.52	329.8	776.0
## 2	Accept	2018.80438	2451.113971	1.2141414	146.00	272.4	604.0
## 3	Enroll	779.97297	929.176190	1.1912928	78.04	118.6	242.0
## 4	Top10perc	27.55856	17.640364	0.6401048	3.00	7.0	15.0
## 5	Top25perc	55.79665	19.804778	0.3549456	18.76	25.8	41.0
## 6	F.Undergrad	3699.90734	4850.420531	1.3109573	324.36	509.8	992.0
## 7	P.Undergrad	855.29858	1522.431887	1.7800005	4.76	20.0	95.0
## 8	Outstate	10440.66924	4023.016484	0.3853217	3737.28	4601.6	7320.0
## 9	Room.Board	4357.52638	1096.696416	0.2516787	2377.60	2735.8	3597.0
## 10	Books	549.38095	165.105360	0.3005298	250.00	350.0	470.0
## 11	Personal	1340.64221	677.071454	0.5050352	400.00	500.0	850.0
## 12	PhD	72.66023	16.328155	0.2247193	28.28	43.8	62.0
## 13	Terminal	79.70270	14.722359	0.1847159	40.04	52.8	71.0
## 14	S.F.Ratio	14.08970	3.958349	0.2809391	5.00	8.3	11.5
## 15	perc.alumni	22.74389	12.391801	0.5448410	3.00	6.0	13.0
## 16	Expend	9660.17117	5221.768440	0.5405462	3869.32	4795.8	6751.0
## 17	Grad.Rate	65.46332	17.177710	0.2624021	23.52	37.0	53.0

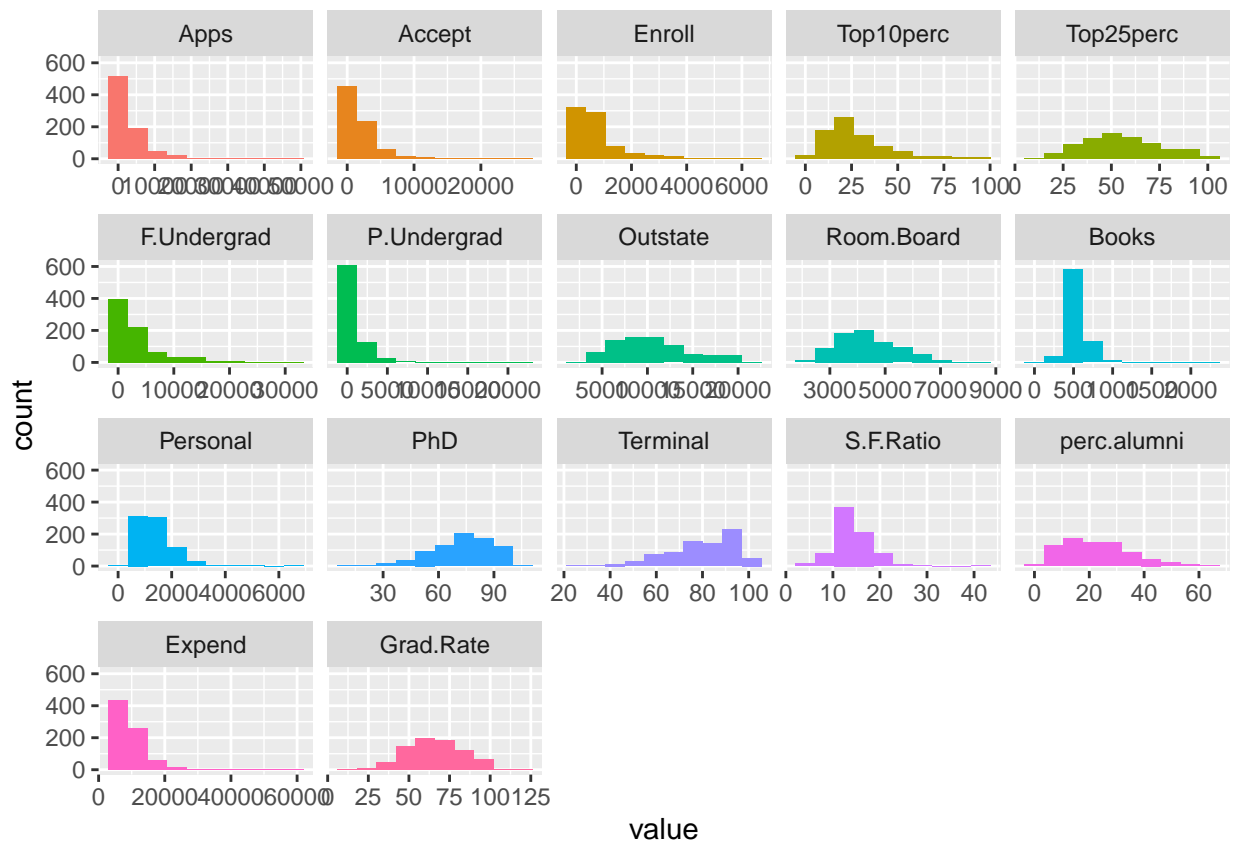
##	p_50	p_75	p_95	p_99	skewness	kurtosis	iqr
## 1	1558.0	3624.0	11066.2	16026.12	3.7165574	29.594559	2848
## 2	1110.0	2424.0	6979.2	11668.08	3.4111259	21.808740	1820
## 3	434.0	902.0	2757.0	4618.84	2.6852679	11.767103	660
## 4	23.0	35.0	65.2	87.48	1.4104871	5.186169	20
## 5	54.0	69.0	93.0	99.00	0.2588394	2.431791	28
## 6	1707.0	4005.0	14477.8	24540.32	2.6054157	10.639436	3013
## 7	353.0	967.0	3303.6	7477.08	5.6813582	57.673296	872
## 8	9990.0	12925.0	18498.0	19677.20	0.5082943	2.581114	5605
## 9	4200.0	5050.0	6382.0	7031.44	0.4764335	2.805940	1453
## 10	500.0	600.0	765.6	1143.00	3.4782933	31.143390	130
## 11	1200.0	1700.0	2488.8	3336.00	1.7391308	10.070544	850
## 12	75.0	85.0	95.0	99.00	-0.7666864	3.553433	23
## 13	82.0	92.0	98.0	100.00	-0.8149652	3.232752	21
## 14	13.6	16.5	21.0	23.72	0.6661462	5.537045	5
## 15	21.0	31.0	46.0	55.00	0.6057190	2.896103	18
## 16	8377.0	10830.0	17974.8	31335.48	3.4526399	21.643210	4079
## 17	65.0	78.0	94.2	100.00	-0.1135575	2.788380	25

##	range_98	range_80
## 1	[192.52, 16026.12]	[457.6, 7675]
## 2	[146, 11668.08]	[361.6, 4814.2]
## 3	[78.04, 4618.84]	[154, 1903.6]
## 4	[3, 87.48]	[10, 50.4]
## 5	[18.76, 99]	[30.6, 85]
## 6	[324.36, 24540.32]	[641, 10024.4]
## 7	[4.76, 7477.08]	[35, 2016.6]
## 8	[3737.28, 19677.2]	[5568.8, 16552.8]
## 9	[2377.6, 7031.44]	[3051.2, 5950]
## 10	[250, 1143]	[400, 700]
## 11	[400, 3336]	[600, 2200]
## 12	[28.28, 99]	[50.6, 92]
## 13	[40.04, 100]	[59, 96]
## 14	[5, 23.72]	[9.9, 19.2]

```
## 15          [3, 55]          [8, 40]
## 16 [3869.32, 31335.48] [5558.2, 14841]
## 17      [23.52, 100]      [44.6, 89]

## Warning: `guides(<scale> = FALSE)` is deprecated. Please use `guides(<scale> =
## "none")` instead.
```





```
## data
##
## 18 Variables      777 Observations
## -----
## Private
##      n missing distinct
##      777      0        2
##
## Value      No  Yes
## Frequency   212 565
## Proportion 0.273 0.727
## -----
## Apps
##      n missing distinct      Info      Mean      Gmd      .05      .10
##      777      0      711        1      3002      3350      329.8      457.6
##      .25      .50      .75      .90      .95
##      776.0    1558.0    3624.0    7675.0    11066.2
##
## lowest :      81    100    141    150    152, highest: 19315 19873 20192 21804 48094
## -----
## Accept
##      n missing distinct      Info      Mean      Gmd      .05      .10
##      777      0      693        1      2019      2127      272.4      361.6
##      .25      .50      .75      .90      .95
##      604.0    1110.0    2424.0    4814.2    6979.2
##
## lowest :      72     90    118    128    130, highest: 13007 13243 15096 18744 26330
```

```

## -----
## Enroll
##      n missing distinct      Info      Mean      Gmd      .05      .10
##      777      0      581      1      780      821.5      118.6      154.0
##      .25      .50      .75      .90      .95
##      242.0      434.0      902.0      1903.6      2757.0
##
## lowest :   35   46   51   55   63, highest: 5705 5873 5874 6180 6392
## -----
## Top10perc
##      n missing distinct      Info      Mean      Gmd      .05      .10
##      777      0      82      0.999      27.56      18.52      7.0      10.0
##      .25      .50      .75      .90      .95
##      15.0      23.0      35.0      50.4      65.2
##
## lowest :   1   2   3   4   5, highest: 87 89 90 95 96
## -----
## Top25perc
##      n missing distinct      Info      Mean      Gmd      .05      .10
##      777      0      89      1      55.8      22.58      25.8      30.6
##      .25      .50      .75      .90      .95
##      41.0      54.0      69.0      85.0      93.0
##
## lowest :    9   12   13   14   16, highest:  96  97  98  99 100
## -----
## F.Undergrad
##      n missing distinct      Info      Mean      Gmd      .05      .10
##      777      0      714      1      3700      4215      509.8      641.0
##      .25      .50      .75      .90      .95
##      992.0      1707.0      4005.0      10024.4      14477.8
##
## lowest :   139   199   201   249   282, highest: 26640 27378 28938 30017 31643
## -----
## P.Undergrad
##      n missing distinct      Info      Mean      Gmd      .05      .10
##      777      0      566      1      855.3      1131      20      35
##      .25      .50      .75      .90      .95
##      95      353      967      2017      3304
##
## lowest :    1    2    3    4    5, highest:  9054  9310 10221 10962 21836
## -----
## Outstate
##      n missing distinct      Info      Mean      Gmd      .05      .10
##      777      0      640      1      10441      4547      4602      5569
##      .25      .50      .75      .90      .95
##      7320      9990      12925      16553      18498
##
## lowest :  2340  2580  2700  3040  3460, highest: 19900 19960 19964 20100 21700
## -----
## Room.Board
##      n missing distinct      Info      Mean      Gmd      .05      .10
##      777      0      553      1      4358      1236      2736      3051
##      .25      .50      .75      .90      .95
##      3597      4200      5050      5950      6382

```

```

##
## lowest : 1780 1880 1920 2146 2190, highest: 7350 7398 7400 7425 8124
## -----
## Books
##      n missing distinct      Info      Mean      Gmd      .05      .10
##      777      0      122    0.982    549.4    152.2    350.0    400.0
##      .25      .50      .75      .90      .95
##      470.0    500.0    600.0    700.0    765.6
##
## lowest :   96  110  120  200  221, highest: 1300 1400 1495 2000 2340
## -----
## Personal
##      n missing distinct      Info      Mean      Gmd      .05      .10
##      777      0      294    0.999    1341    710.3     500     600
##      .25      .50      .75      .90      .95
##      850    1200    1700    2200    2489
##
## lowest :  250  300  350  400  420, highest: 4110 4200 4288 4913 6800
## -----
## PhD
##      n missing distinct      Info      Mean      Gmd      .05      .10
##      777      0       78       1    72.66    18.18     43.8     50.6
##      .25      .50      .75      .90      .95
##      62.0    75.0    85.0    92.0    95.0
##
## lowest :   8  10  14  16  22, highest:  97  98  99 100 103
## -----
## Terminal
##      n missing distinct      Info      Mean      Gmd      .05      .10
##      777      0       65    0.999    79.7    16.37     52.8     59.0
##      .25      .50      .75      .90      .95
##      71.0    82.0    92.0    96.0    98.0
##
## lowest :  24  25  30  33  35, highest:  96  97  98  99 100
## -----
## S.F.Ratio
##      n missing distinct      Info      Mean      Gmd      .05      .10
##      777      0      173       1    14.09    4.325     8.3     9.9
##      .25      .50      .75      .90      .95
##      11.5    13.6    16.5    19.2    21.0
##
## lowest :  2.5  2.9  3.3  3.9  4.3, highest: 27.2 27.6 27.8 28.8 39.8
## -----
## perc.alumni
##      n missing distinct      Info      Mean      Gmd      .05      .10
##      777      0       61    0.999    22.74    13.95       6       8
##      .25      .50      .75      .90      .95
##      13      21      31      40      46
##
## lowest :  0  1  2  3  4, highest: 57 58 60 63 64
## -----
## Expend
##      n missing distinct      Info      Mean      Gmd      .05      .10
##      777      0      744       1    9660    4650    4796    5558

```

```
##      .25      .50      .75      .90      .95
##    6751    8377   10830   14841   17975
##
## lowest :   3186   3365   3480   3605   3733, highest: 40386 41766 42926 45702 56233
## -----
## Grad.Rate
##      n missing distinct      Info      Mean      Gmd      .05      .10
##    777      0       81        1    65.46    19.48    37.0    44.6
##      .25      .50      .75      .90      .95
##    53.0    65.0    78.0    89.0    94.2
##
## lowest :   10   15   18   21   22, highest:   97   98   99 100 118
## -----
```

Split the data set into a training set and a test set.

```
College_split <- initial_split(College, strata = Apps, prop = 0.5)
College_split
```

```
## <Analysis/Assess/Total>
## <388/389/777>
College_train <- training(College_split)
College_test  <- testing(College_split)
```

Fit a linear model using least squares on the training set, and report the test error obtained.

```
lm_spec <- linear_reg() %>%
  set_mode("regression") %>%
  set_engine("lm")

lm_recipe <-
  recipe(formula = Apps ~ ., data = College_train) %>%
  step_dummy(all_nominal_predictors()) %>%
  step_normalize(all_numeric_predictors())
```

Note: We use the variables `Accept` and `Enroll` as independent variables here. These variables describe the number of accepted applicants and the number of new enrolled students, respectively. Depending on the application one might not know these variables when predicting the the number of applications.

```
lm_workflow <- workflow() %>%
  add_recipe(lm_recipe) %>%
  add_model(lm_spec)

lm_fit <- lm_workflow %>% fit(College_train)

augment(lm_fit, new_data = College_test) %>%
  select(Apps, .pred)
```

```
## # A tibble: 389 x 2
##   Apps .pred
##   <dbl> <dbl>
## 1 1660 1349.
## 2  587  625.
## 3 1038  984.
## 4  582  574.
## 5 1179 1731.
## 6 1420 1211.
```



```
## 7 1130 463.
## 8 3540 3107.
## 9 619 390.
## 10 12809 15205.
## # ... with 379 more rows

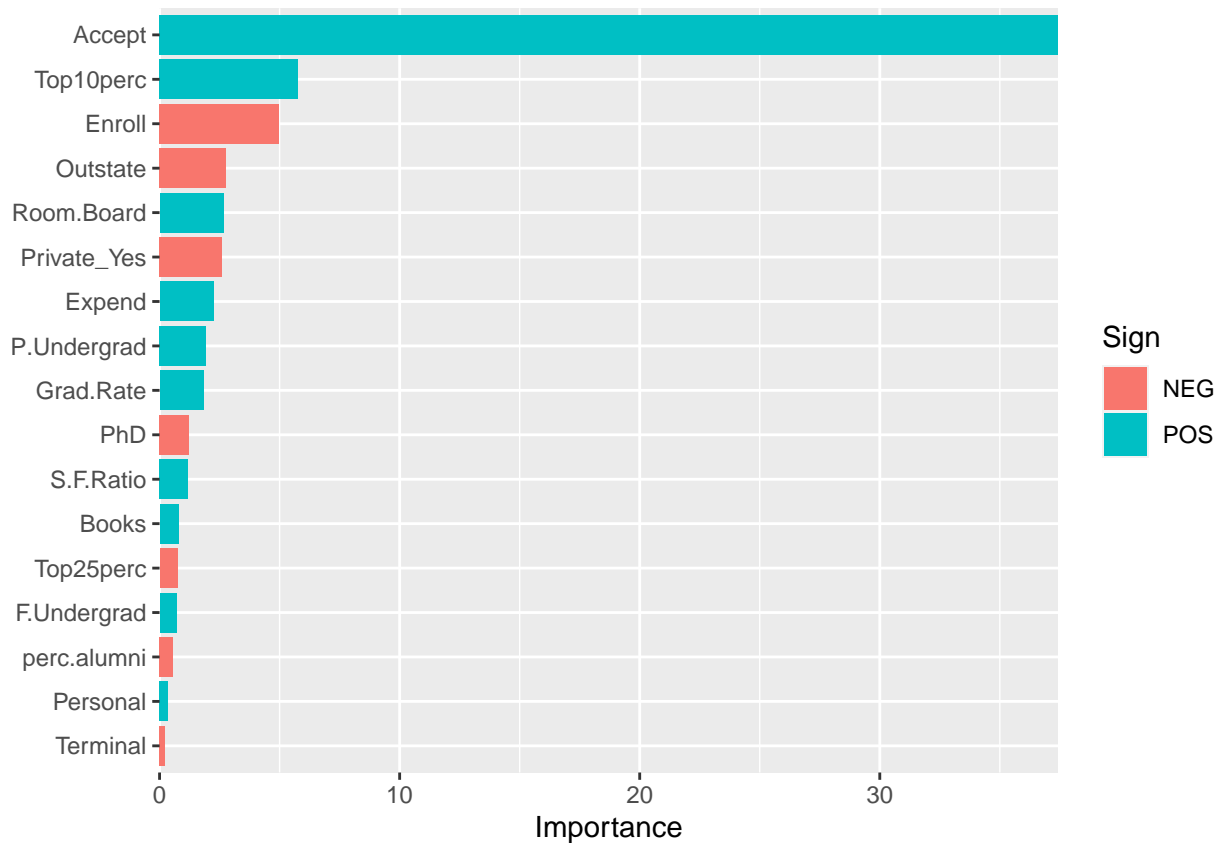
augment(lm_fit, new_data = College_test) %>%
  rmse(truth = Apps, estimate = .pred)

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard     1183.
```

The rmse of the test error is 1183.

Out of curiosity, let's look at the feature importance.

```
lm_fit %>%
  extract_fit_parsnip() %>%
  vi(lambda = best_penalty$penalty) %>%
  mutate(
    Importance = abs(Importance),
    Variable = fct_reorder(Variable, Importance)
  ) %>%
  ggplot(aes(x = Importance, y = Variable, fill = Sign)) +
  geom_col() +
  scale_x_continuous(expand = c(0, 0)) +
  labs(y = NULL)
```



Fit a ridge regression model on the training set, with λ chosen by cross-validation. Report the test error obtained.

```
ridge_spec <- linear_reg(mixture = 0, penalty = tune()) %>%
  set_mode("regression") %>%
  set_engine("glmnet")

ridge_workflow <- workflow() %>%
  add_recipe(recipe = lm_recipe) %>%
  add_model(ridge_spec)

College_fold <- vfold_cv(College_train, v = 10)
College_fold

## # 10-fold cross-validation
## # A tibble: 10 x 2
##   splits      id
##   <list>    <chr>
## 1 <split [349/39]> Fold01
## 2 <split [349/39]> Fold02
## 3 <split [349/39]> Fold03
## 4 <split [349/39]> Fold04
## 5 <split [349/39]> Fold05
## 6 <split [349/39]> Fold06
## 7 <split [349/39]> Fold07
## 8 <split [349/39]> Fold08
## 9 <split [350/38]> Fold09
## 10 <split [350/38]> Fold10

penalty_grid <- grid_regular(
  penalty(range = c(-5, 5)), # penalty automatically uses log scale
  levels = 50
)
penalty_grid

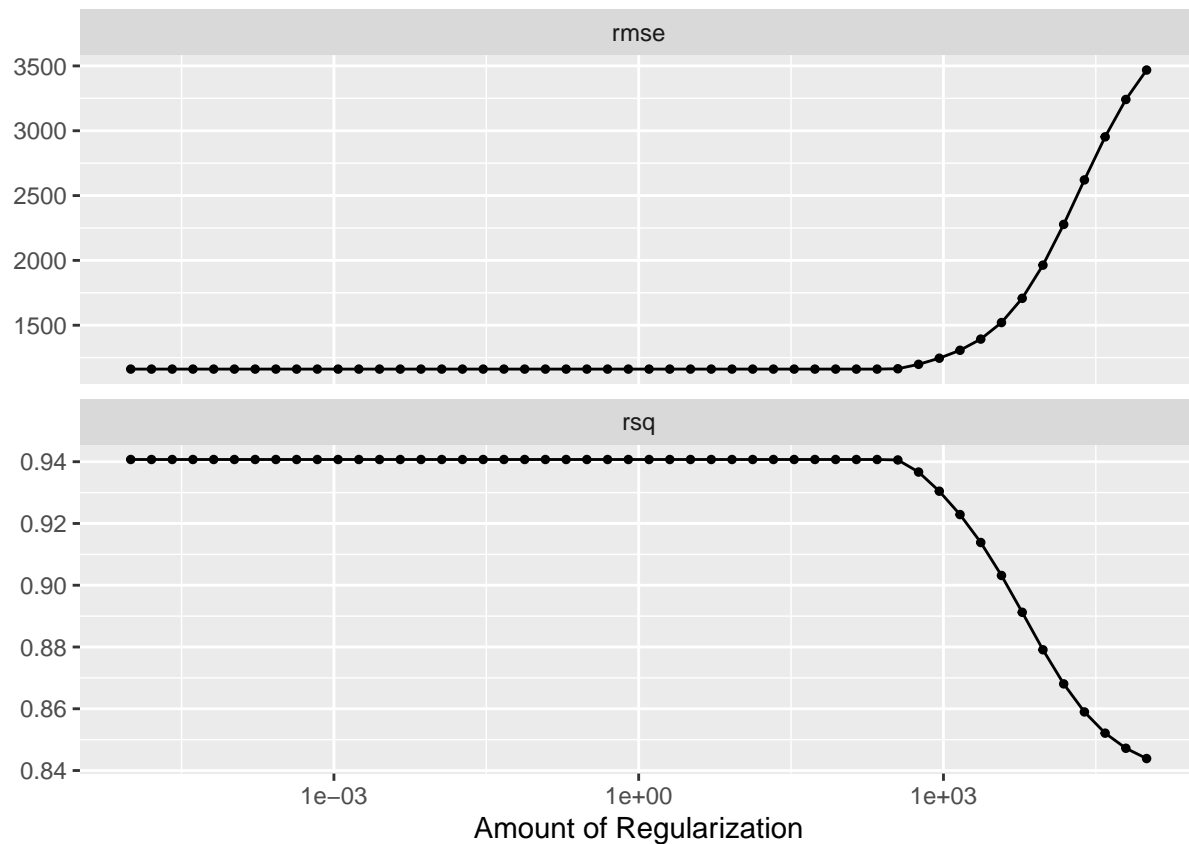
## # A tibble: 50 x 1
##   penalty
##   <dbl>
## 1 0.00001
## 2 0.0000160
## 3 0.0000256
## 4 0.0000409
## 5 0.0000655
## 6 0.000105
## 7 0.000168
## 8 0.000268
## 9 0.000429
## 10 0.000687
## # ... with 40 more rows

tune_res <- tune_grid(
  ridge_workflow,
  resamples = College_fold,
  grid = penalty_grid
)
```

```
tune_res
```

```
## # Tuning results
## # 10-fold cross-validation
## # A tibble: 10 x 4
##   splits      id      .metrics      .notes
##   <list>    <chr> <list>      <list>
## 1 <split [349/39]> Fold01 <tibble [100 x 5]> <tibble [0 x 3]>
## 2 <split [349/39]> Fold02 <tibble [100 x 5]> <tibble [0 x 3]>
## 3 <split [349/39]> Fold03 <tibble [100 x 5]> <tibble [0 x 3]>
## 4 <split [349/39]> Fold04 <tibble [100 x 5]> <tibble [0 x 3]>
## 5 <split [349/39]> Fold05 <tibble [100 x 5]> <tibble [0 x 3]>
## 6 <split [349/39]> Fold06 <tibble [100 x 5]> <tibble [0 x 3]>
## 7 <split [349/39]> Fold07 <tibble [100 x 5]> <tibble [0 x 3]>
## 8 <split [349/39]> Fold08 <tibble [100 x 5]> <tibble [0 x 3]>
## 9 <split [350/38]> Fold09 <tibble [100 x 5]> <tibble [0 x 3]>
## 10 <split [350/38]> Fold10 <tibble [100 x 5]> <tibble [0 x 3]>
```

```
autoplot(tune_res)
```



```
best_penalty <- select_best(tune_res, metric = "rmse")
best_penalty
```

```
## # A tibble: 1 x 2
##   penalty .config
##   <dbl> <chr>
## 1 0.00001 Preprocessor1_Model01
```

```
ridge_final <- finalize_workflow(ridge_workflow, best_penalty)
ridge_final_fit <- ridge_final %>% fit(College_train)
```

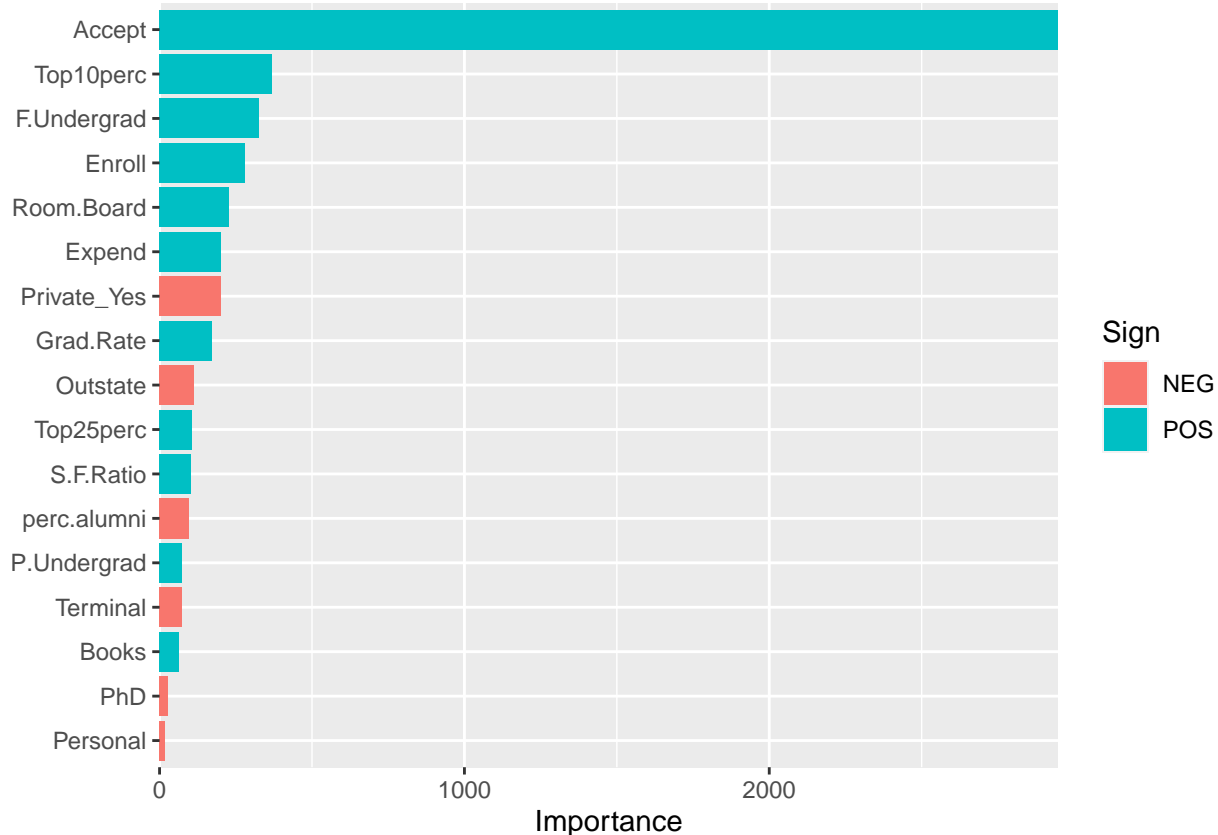
```
augment(ridge_final_fit, new_data = College_test) %>%
  rmse(truth = Apps, estimate = .pred)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard      1178.
```

The test error is 1178, which is slightly lower than a linear model without regularization.

Again let's also look at the feature importance.

```
ridge_final_fit %>%
  extract_fit_parsnip() %>%
  vi(lambda = best_penalty$penalty) %>%
  mutate(
    Importance = abs(Importance),
    Variable = fct_reorder(Variable, Importance)
  ) %>%
  ggplot(aes(x = Importance, y = Variable, fill = Sign)) +
  geom_col() +
  scale_x_continuous(expand = c(0, 0)) +
  labs(y = NULL)
```



Fit a lasso model on the training set, with λ chosen by cross-validation. Report the test error obtained, along with the number of non-zero coefficient estimates.

```

lasso_spec <- linear_reg(mixture = 1, penalty = tune()) %>%
  set_mode("regression") %>%
  set_engine("glmnet")

lasso_workflow <- workflow() %>%
  add_recipe(recipe = lm_recipe) %>%
  add_model(lasso_spec)

penalty_grid <- grid_regular(
  penalty(range = c(-5, 2)),
  levels = 50
)

tune_res <- tune_grid(
  lasso_workflow,
  resamples = College_fold,
  grid = penalty_grid
)
tune_res

## # Tuning results
## # 10-fold cross-validation
## # A tibble: 10 x 4
##   splits          id      .metrics      .notes
##   <list>         <chr>    <list>      <list>
## 1 <split [349/39]> Fold01 <tibble [100 x 5]> <tibble [0 x 3]>
## 2 <split [349/39]> Fold02 <tibble [100 x 5]> <tibble [0 x 3]>
## 3 <split [349/39]> Fold03 <tibble [100 x 5]> <tibble [0 x 3]>
## 4 <split [349/39]> Fold04 <tibble [100 x 5]> <tibble [0 x 3]>
## 5 <split [349/39]> Fold05 <tibble [100 x 5]> <tibble [0 x 3]>
## 6 <split [349/39]> Fold06 <tibble [100 x 5]> <tibble [0 x 3]>
## 7 <split [349/39]> Fold07 <tibble [100 x 5]> <tibble [0 x 3]>
## 8 <split [349/39]> Fold08 <tibble [100 x 5]> <tibble [0 x 3]>
## 9 <split [350/38]> Fold09 <tibble [100 x 5]> <tibble [0 x 3]>
## 10 <split [350/38]> Fold10 <tibble [100 x 5]> <tibble [0 x 3]>

best_penalty <- select_best(tune_res, metric = "rmse")
best_penalty

## # A tibble: 1 x 2
##   penalty .config
##   <dbl> <chr>
## 1    26.8 Preprocessor1_Model46

lasso_final <- finalize_workflow(lasso_workflow, best_penalty)

lasso_final_fit <- fit(lasso_final, data = College_train)

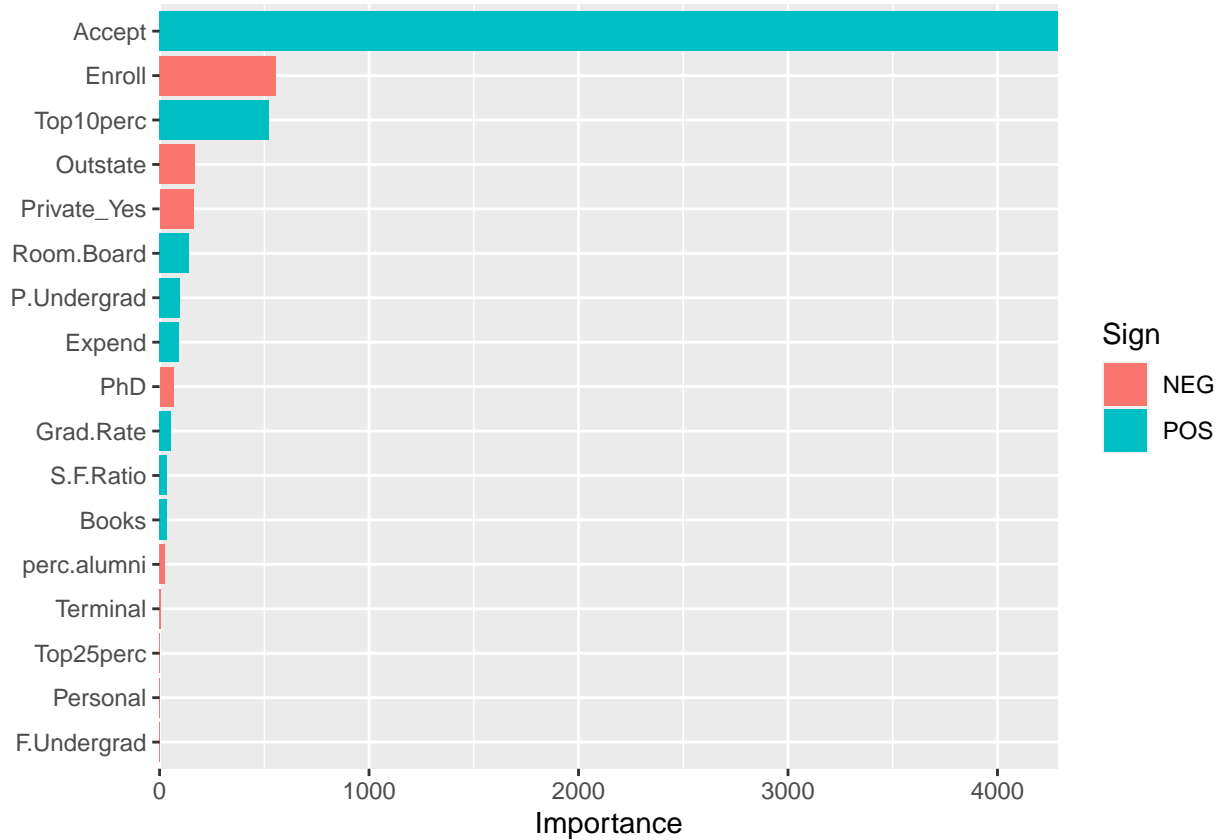
augment(ridge_final_fit, new_data = College_test) %>%
  rmse(truth = Apps, estimate = .pred)

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>         <dbl>
## 1 rmse   standard       1178.

```

The rmse is 1178, which is the same as for Ridge Regression.

```
lasso_final_fit %>%
  extract_fit_parsnip() %>%
  vi(lambda = best_penalty$penalty) %>%
  mutate(
    Importance = abs(Importance),
    Variable = fct_reorder(Variable, Importance)
  ) %>%
  ggplot(aes(x = Importance, y = Variable, fill = Sign)) +
  geom_col() +
  scale_x_continuous(expand = c(0, 0)) +
  labs(y = NULL)
```



Exercise 2 Fit some of the non-linear models (polynomial regression, splines) discussed in the lecture to the Auto data set. Is there evidence for non-linear relationships in this data set? Create some informative plots to justify your answer.

```
# Recipe
lasso_pl_recipe <-
  recipe(formula = Apps ~ ., data = College_train) %>%
  step_poly(all_numeric_predictors(), degree = 2, options = list(raw = TRUE)) %>%
  step_novel(all_nominal_predictors()) %>%
  step_dummy(all_nominal_predictors()) %>%
  step_zv(all_predictors()) %>%
  step_normalize(all_predictors())

# Specification
lasso_pl_spec <-
```

```

linear_reg(mixture = 1, penalty = tune()) %>%
  set_mode("regression") %>%
  set_engine("glmnet")

# Workflow
lasso_pl_workflow <- workflow() %>%
  add_recipe(lasso_pl_recipe) %>%
  add_model(lasso_pl_spec)

# Grid
penalty_grid <- grid_regular(
  penalty(range = c(-5, 2)),
  levels = 50
)

# Tune model
tune_res <- tune_grid(
  lasso_pl_workflow,
  resamples = College_fold,
  grid = penalty_grid
)

# Finalize model
lasso_pl_final <- finalize_workflow(lasso_pl_workflow, best_penalty)
lasso_pl_final_fit <- fit(lasso_pl_final, data = College_train)

# Check RMSE
augment(lasso_pl_final_fit, new_data = College_test) %>%
  rmse(truth = Apps, estimate = .pred)

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard      1083.

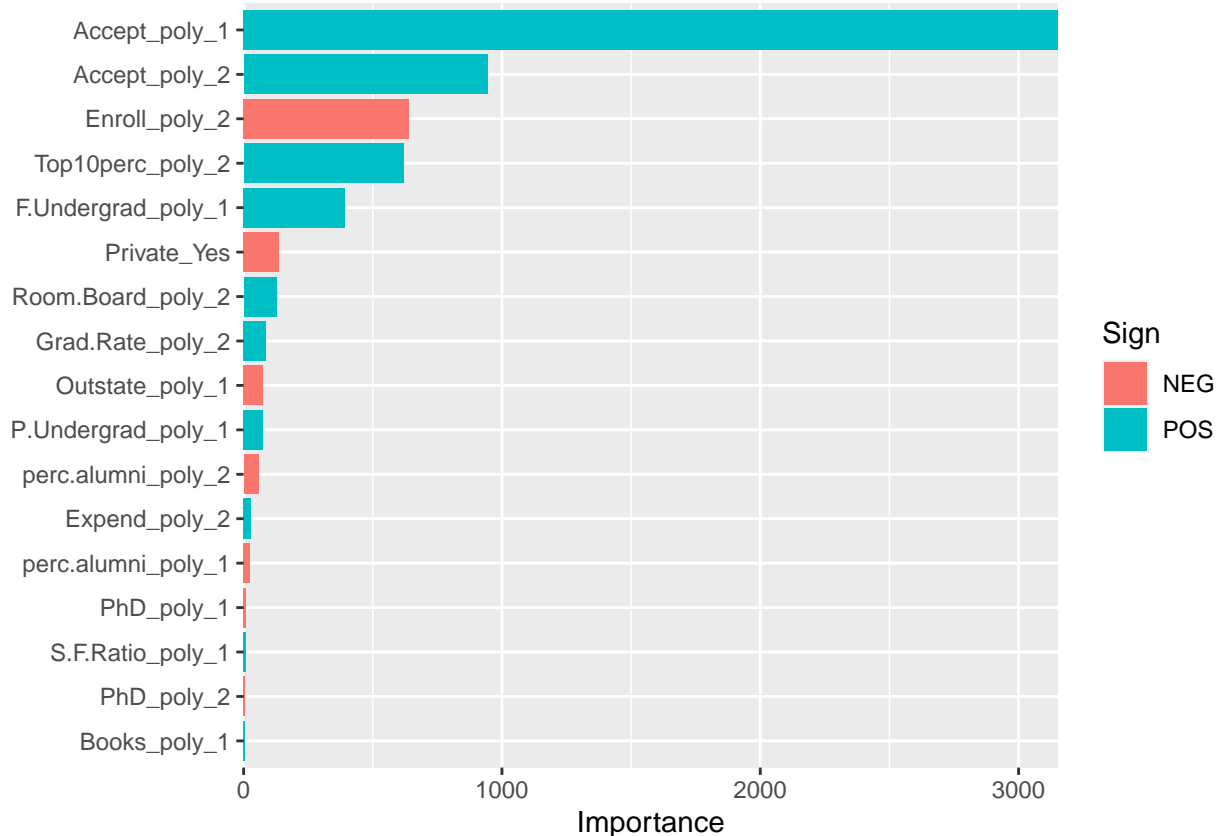
```

Using polynomial features reduced the error by 100, which is almost 10%. This indicates that there are non-linear relationships in the data. Let's look at the feature importance to see which polynomials are important.

```

lasso_pl_final_fit %>%
  extract_fit_parsnip() %>%
  vi(lambda = best_penalty$penalty) %>%
  filter(Importance > 2) %>%
  mutate(
    Importance = abs(Importance),
    Variable = fct_reorder(Variable, Importance)
  ) %>%
  ggplot(aes(x = Importance, y = Variable, fill = Sign)) +
  geom_col() +
  scale_x_continuous(expand = c(0, 0)) +
  labs(y = NULL)

```



As we can see, not surprisingly **Accept** is still the most important feature together with its polynomial. Let's run another polynomial (lasso) regression just with **Accept** so that we can visualize the relationship.

```
# Recipe
lasso_pl_recipe <-
  recipe(formula = Apps ~ Accept, data = College_train) %>%
  step_poly(all_numeric_predictors(), degree = 2, options = list(raw = TRUE)) %>%
  step_normalize(all_predictors())

# Specification
lasso_pl_spec <-
  linear_reg(mixture = 1, penalty = tune()) %>%
  set_mode("regression") %>%
  set_engine("glmnet")

# Workflow
lasso_pl_workflow <- workflow() %>%
  add_recipe(lasso_pl_recipe) %>%
  add_model(lasso_pl_spec)

# Grid
penalty_grid <- grid_regular(
  penalty(range = c(-5, 2)),
  levels = 50
)

# Tune model
tune_res <- tune_grid(
```



```

lasso_pl_workflow,
  resamples = College_fold,
  grid = penalty_grid
)

# Finalize model
lasso_pl_final <- finalize_workflow(lasso_pl_workflow, best_penalty)
lasso_pl_final_fit <- fit(lasso_pl_final, data = College_train)

# Check RMSE
augment(lasso_pl_final_fit, new_data = College_test) %>%
  rmse(truth = Apps, estimate = .pred)

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard      1439.

accept_min <- min(College$Accept)
accept_max <- max(College$Accept)
accept_range <- tibble(Accept = seq(accept_min, accept_max))
accept_range

## # A tibble: 26,259 x 1
##   Accept
##   <int>
## 1     72
## 2     73
## 3     74
## 4     75
## 5     76
## 6     77
## 7     78
## 8     79
## 9     80
## 10    81
## # ... with 26,249 more rows

regression_lines <- bind_cols(
  predict(lasso_pl_final_fit, new_data = accept_range),
  accept_range
)

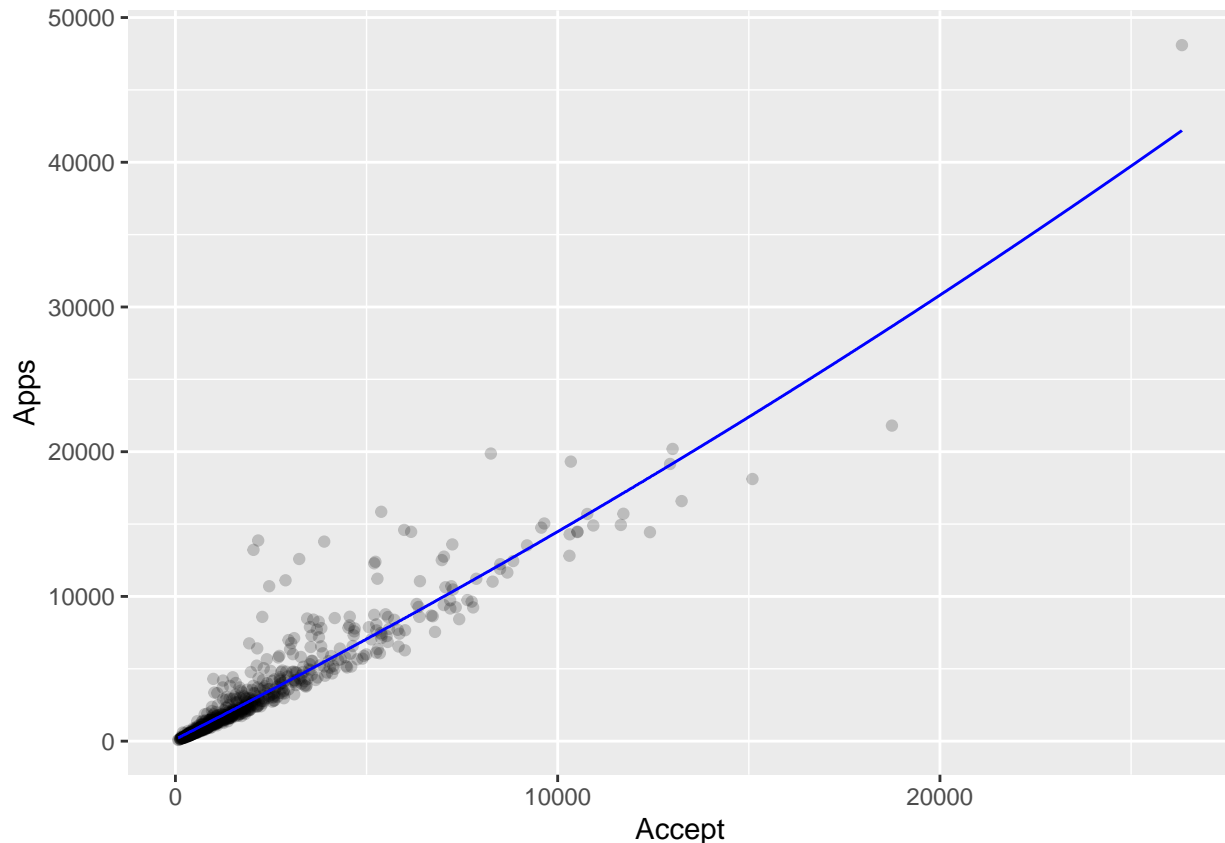
College

## # A tibble: 777 x 18
##   Private Apps Accept Enroll Top10perc Top25perc F.Undergrad P.Undergrad
##   <fct>   <dbl> <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 Yes    1660  1232   721     23     52    2885    537
## 2 Yes    2186  1924   512     16     29    2683   1227
## 3 Yes    1428  1097   336     22     50    1036    99
## 4 Yes     417   349   137     60     89     510    63
## 5 Yes     193   146    55     16     44     249   869
## 6 Yes     587   479   158     38     62     678    41
## 7 Yes     353   340   103     17     45     416   230

```

```
## 8 Yes      1899  1720   489      37      68      1594      32
## 9 Yes      1038   839   227      30      63       973     306
## 10 Yes       582   498   172      21      44       799      78
## # ... with 767 more rows, and 10 more variables: Outstate <dbl>,
## #   Room.Board <dbl>, Books <dbl>, Personal <dbl>, PhD <dbl>, Terminal <dbl>,
## #   S.F.Ratio <dbl>, perc.alumni <dbl>, Expend <dbl>, Grad.Rate <dbl>
```

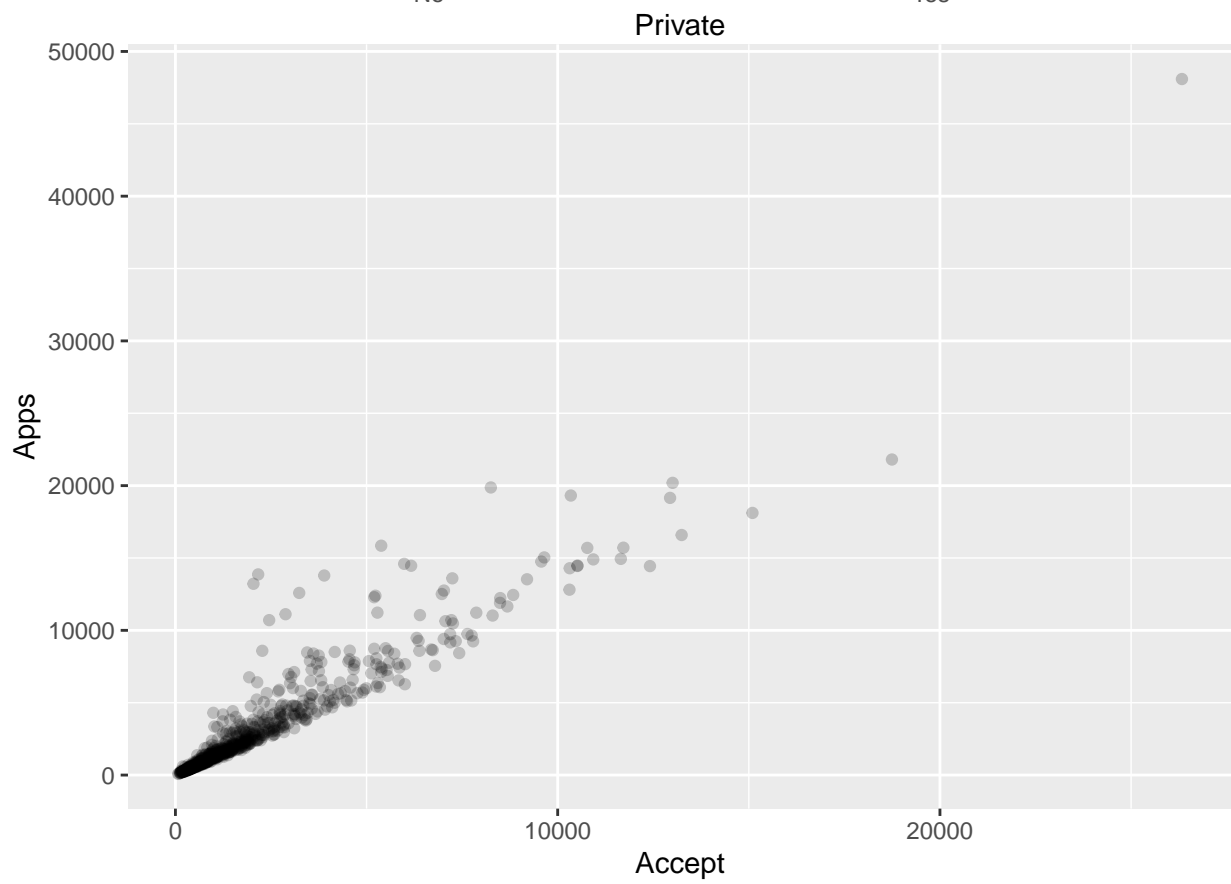
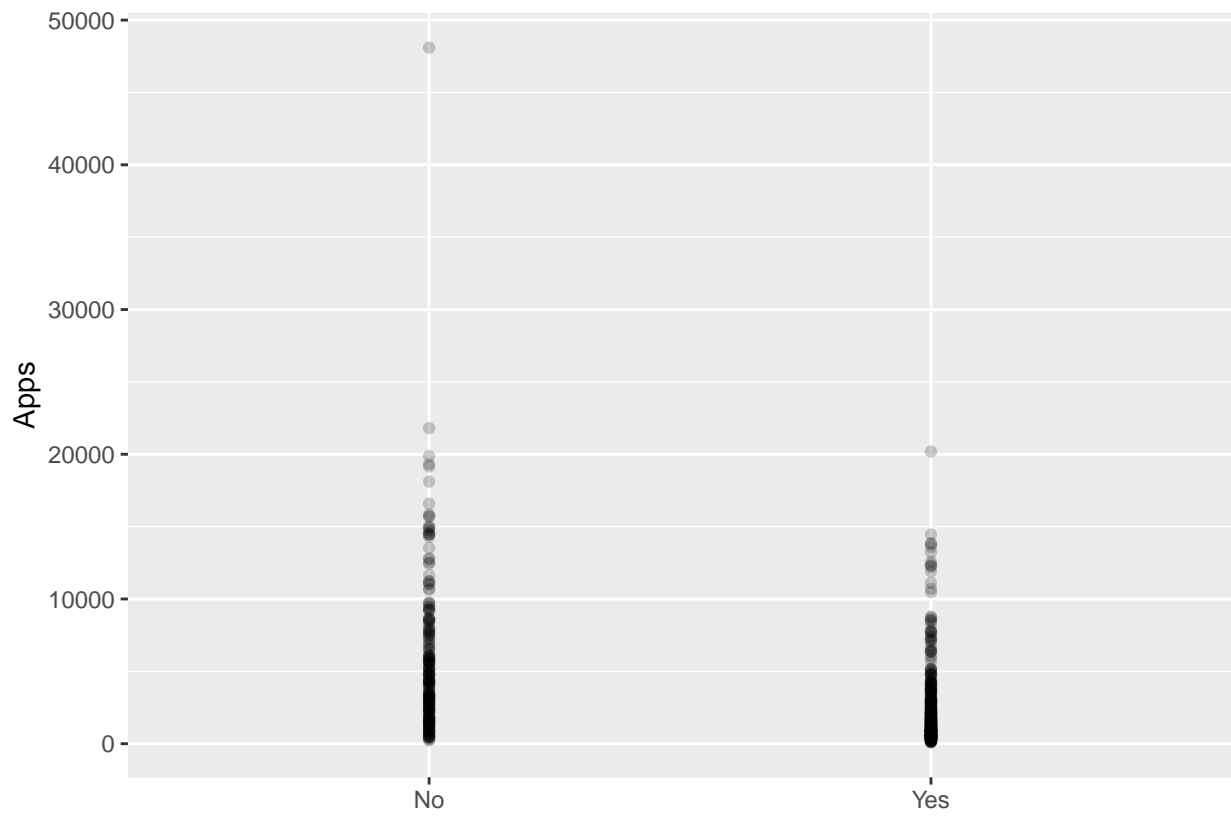
```
College %>%
  ggplot(aes(Accept, Apps)) +
  geom_point(alpha = 0.2) +
  geom_line(aes(y = .pred), data = regression_lines, color = "blue")
```

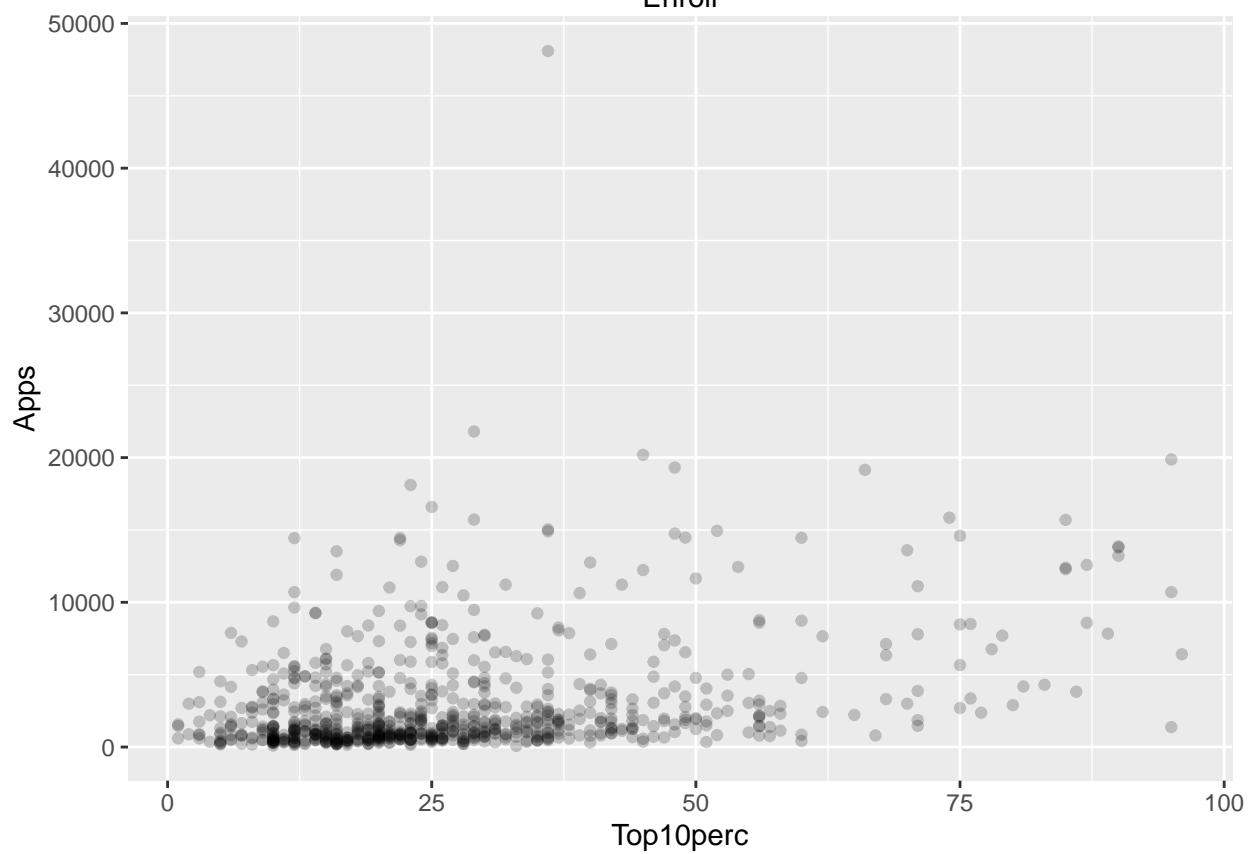
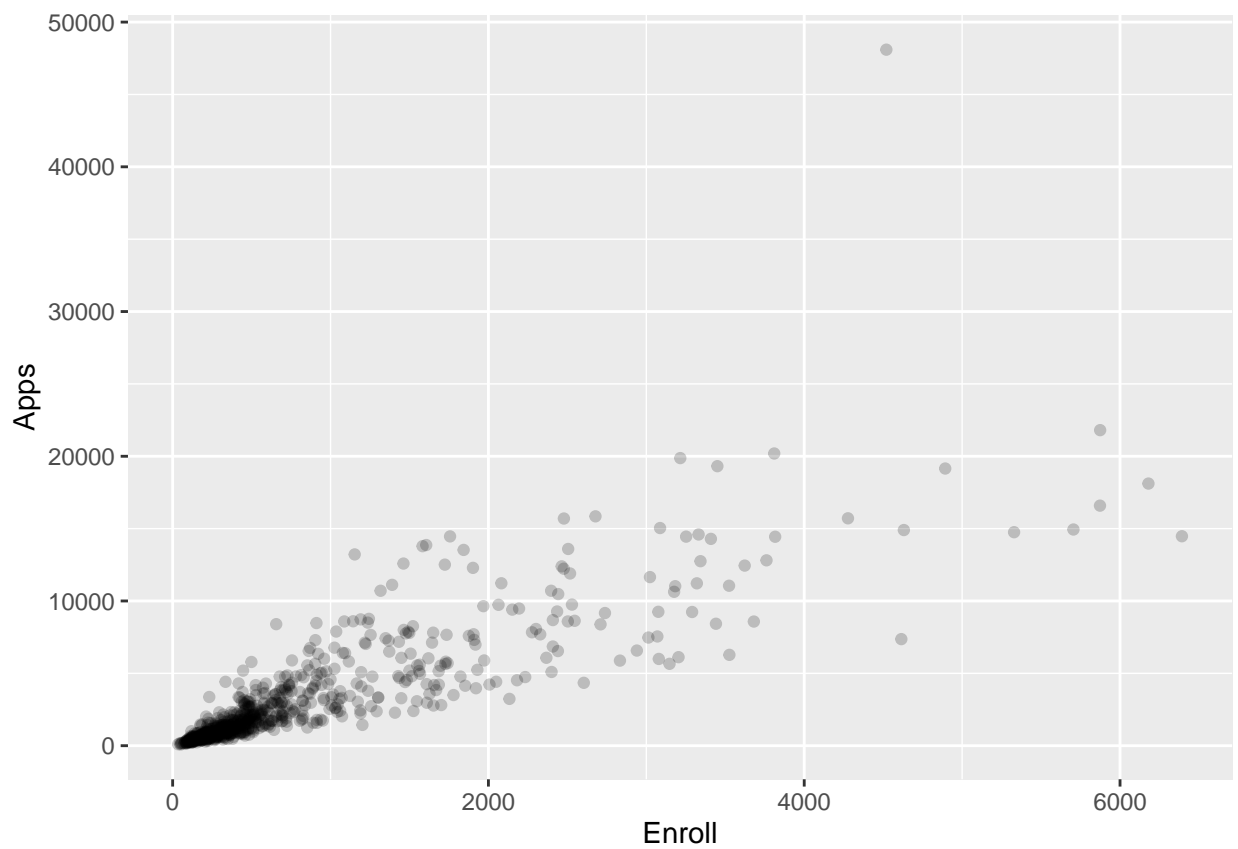


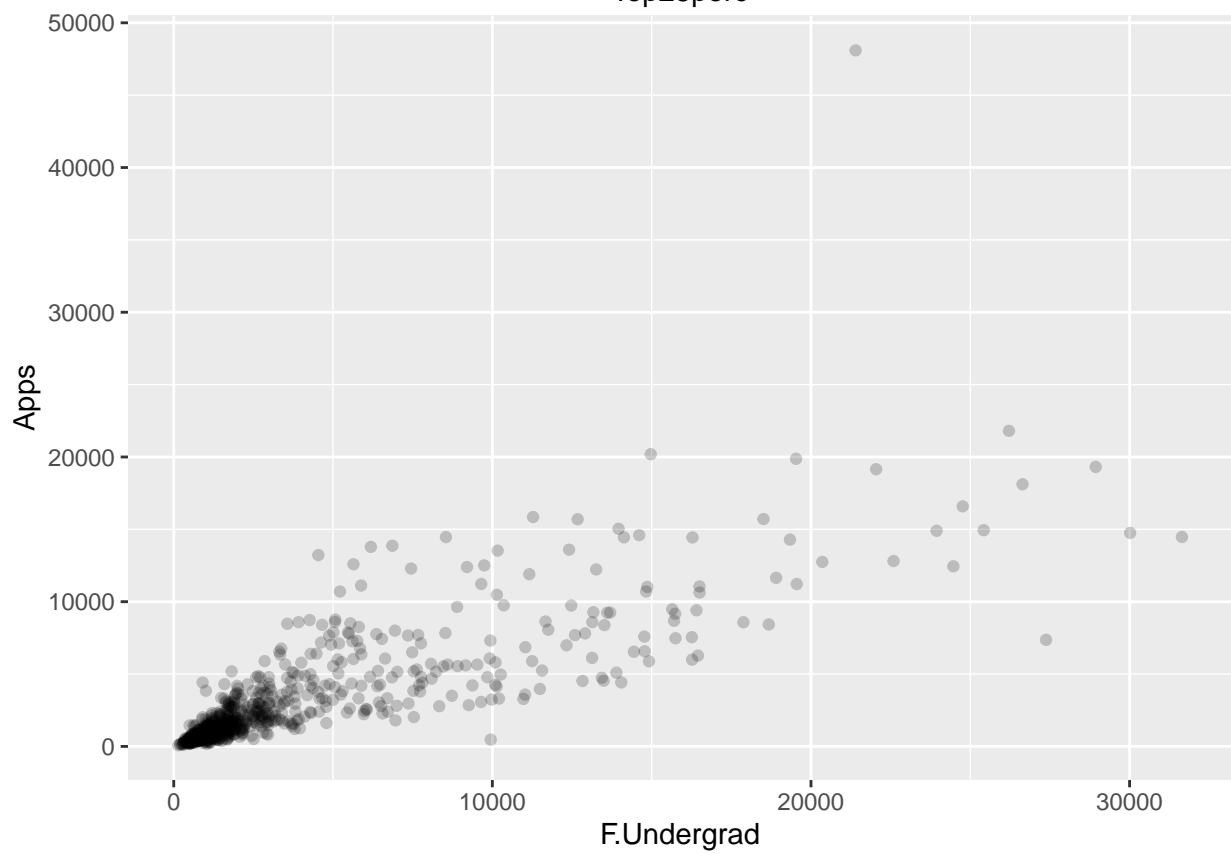
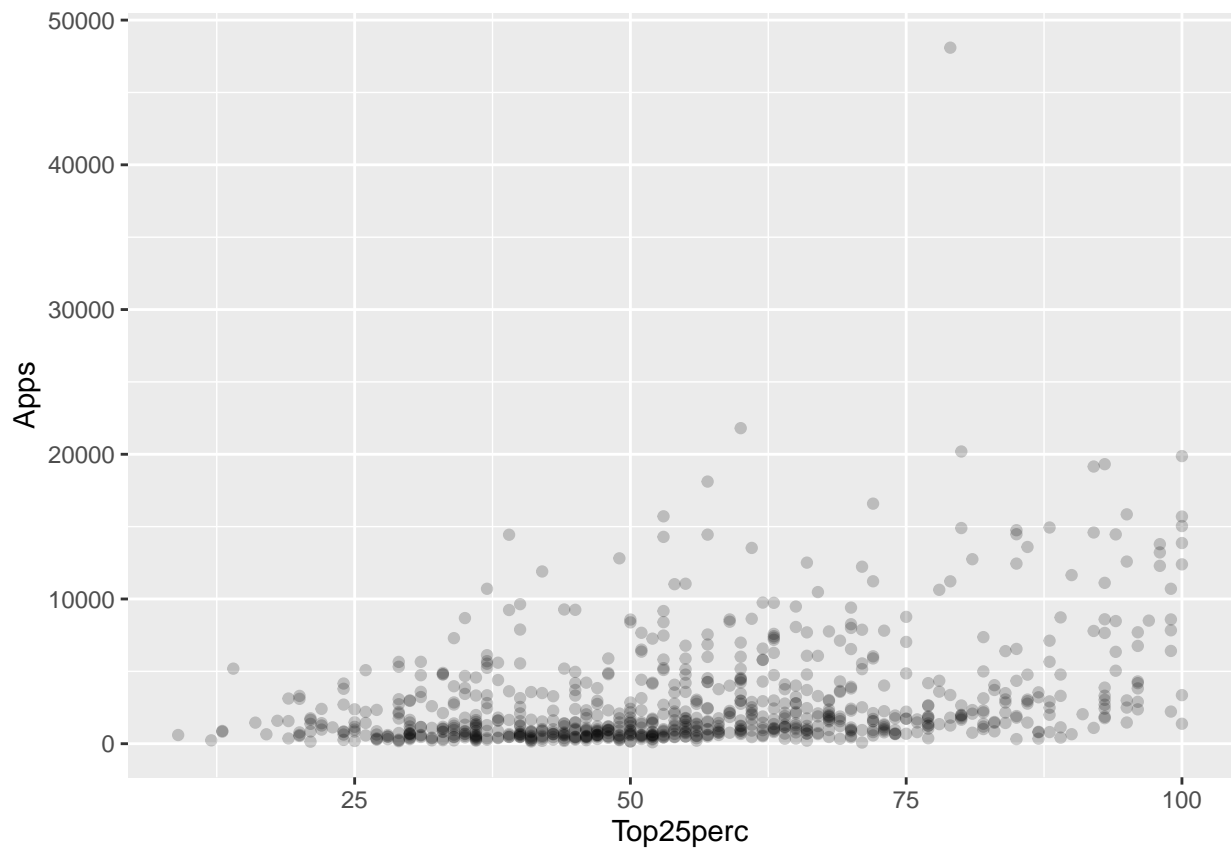
Ok this does look pretty linear. Let us take a look at the scatter plot between each variable and Accept instead.

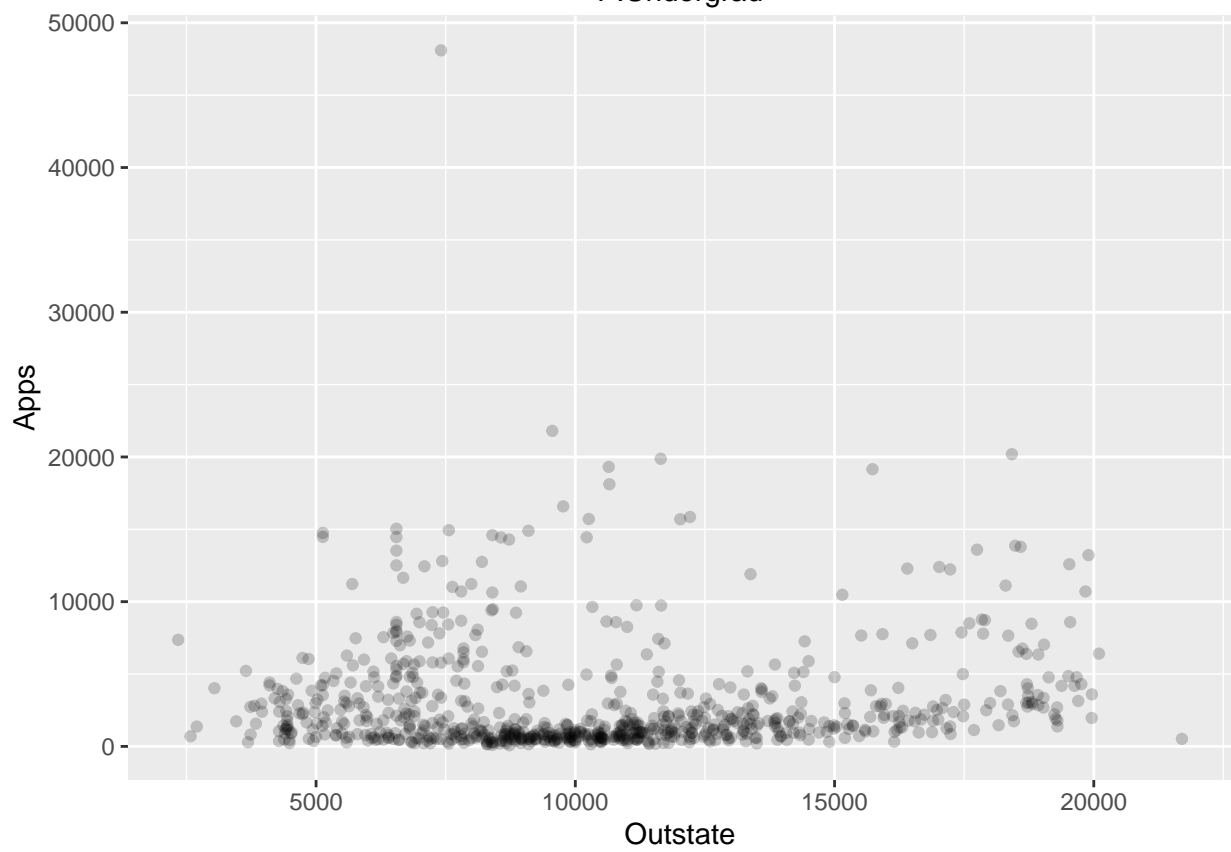
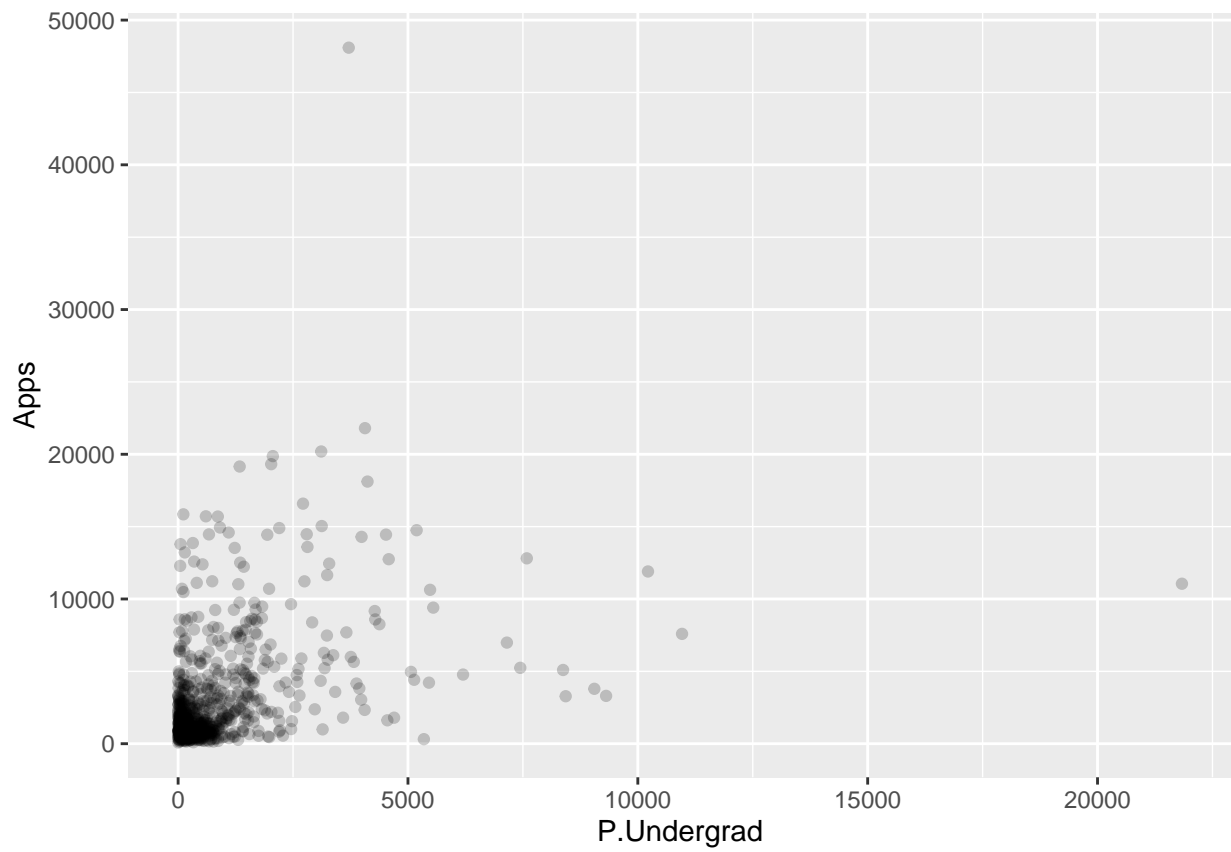
```
# plot scatter of all variabes in College vs Apps
# get all colnames except for Apps
colnames <- colnames(College)
colnames <- colnames[colnames != "Apps"]

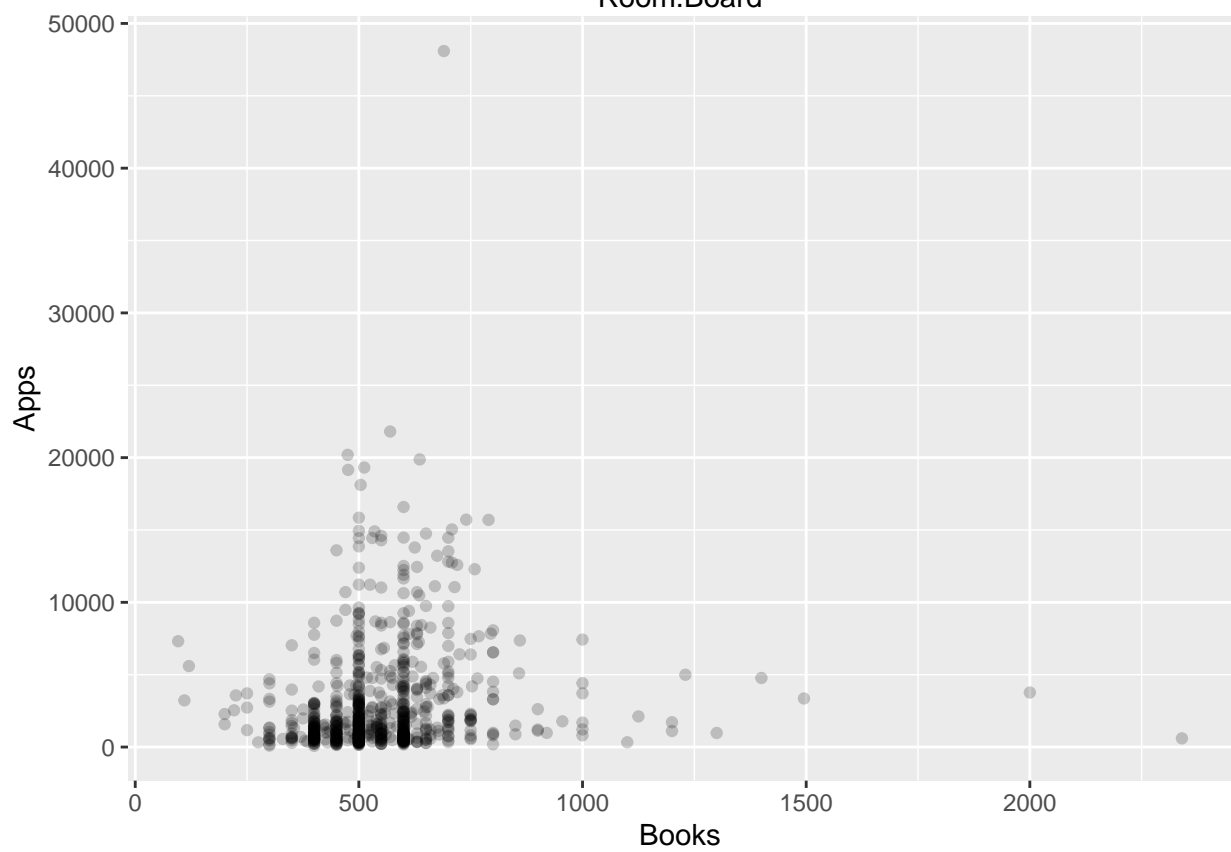
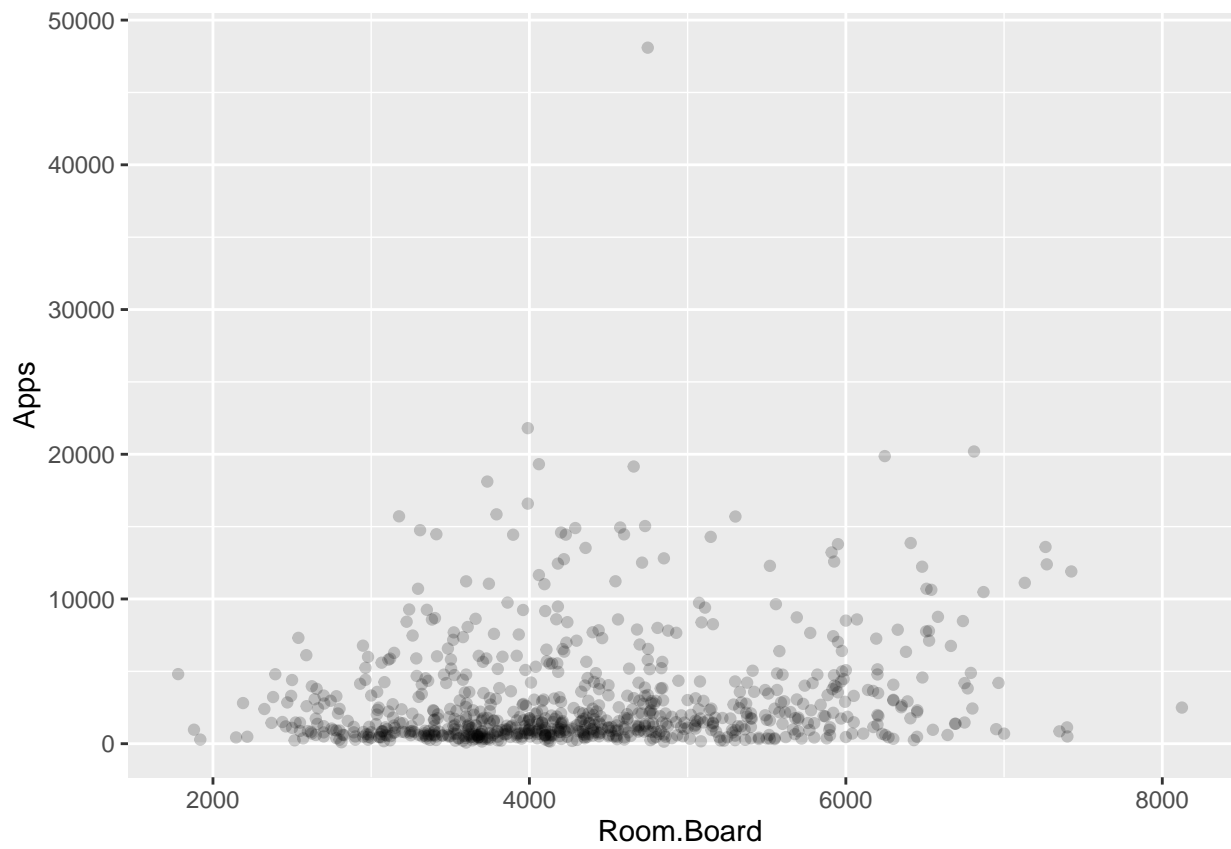
for (colname in colnames) {
  pl <- College %>%
    ggplot(aes_string(x = colname, y = "Apps")) +
    geom_point(alpha = 0.2)
  print(pl)
}
```

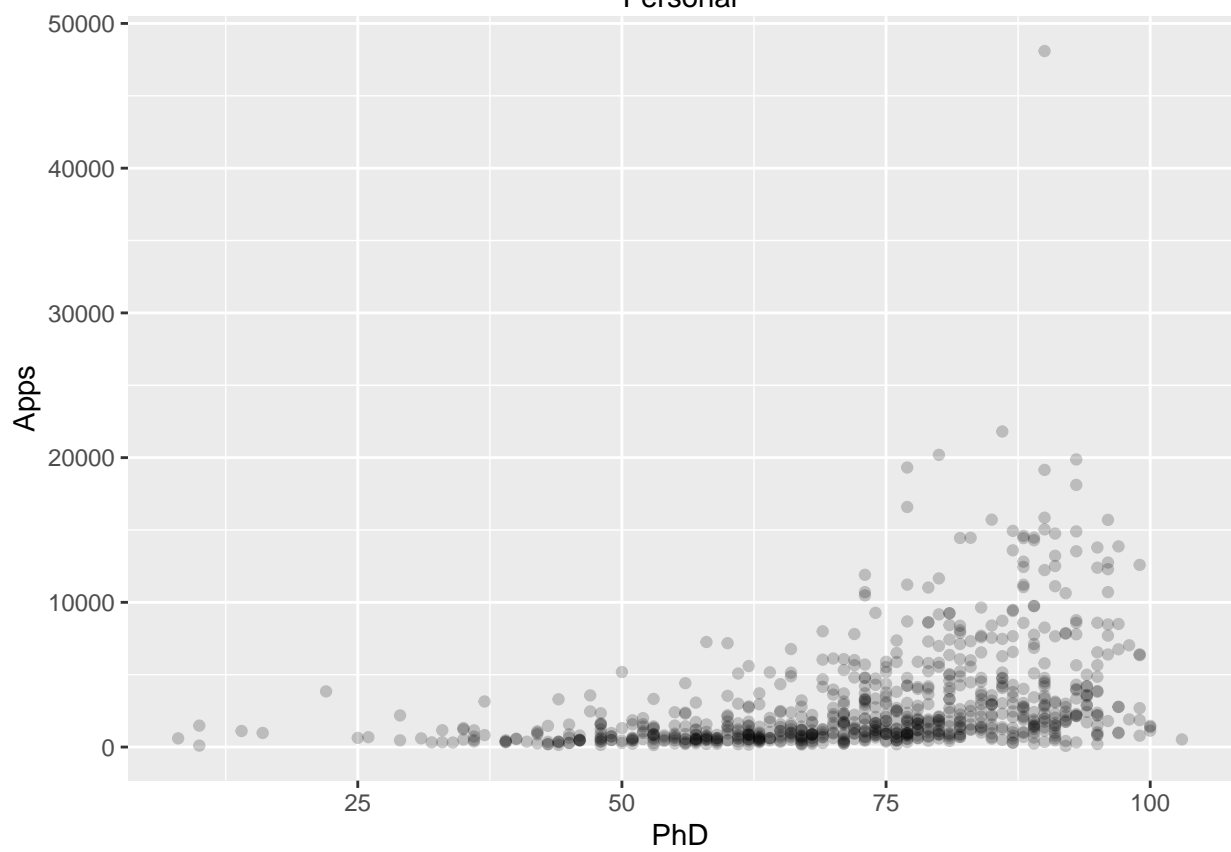
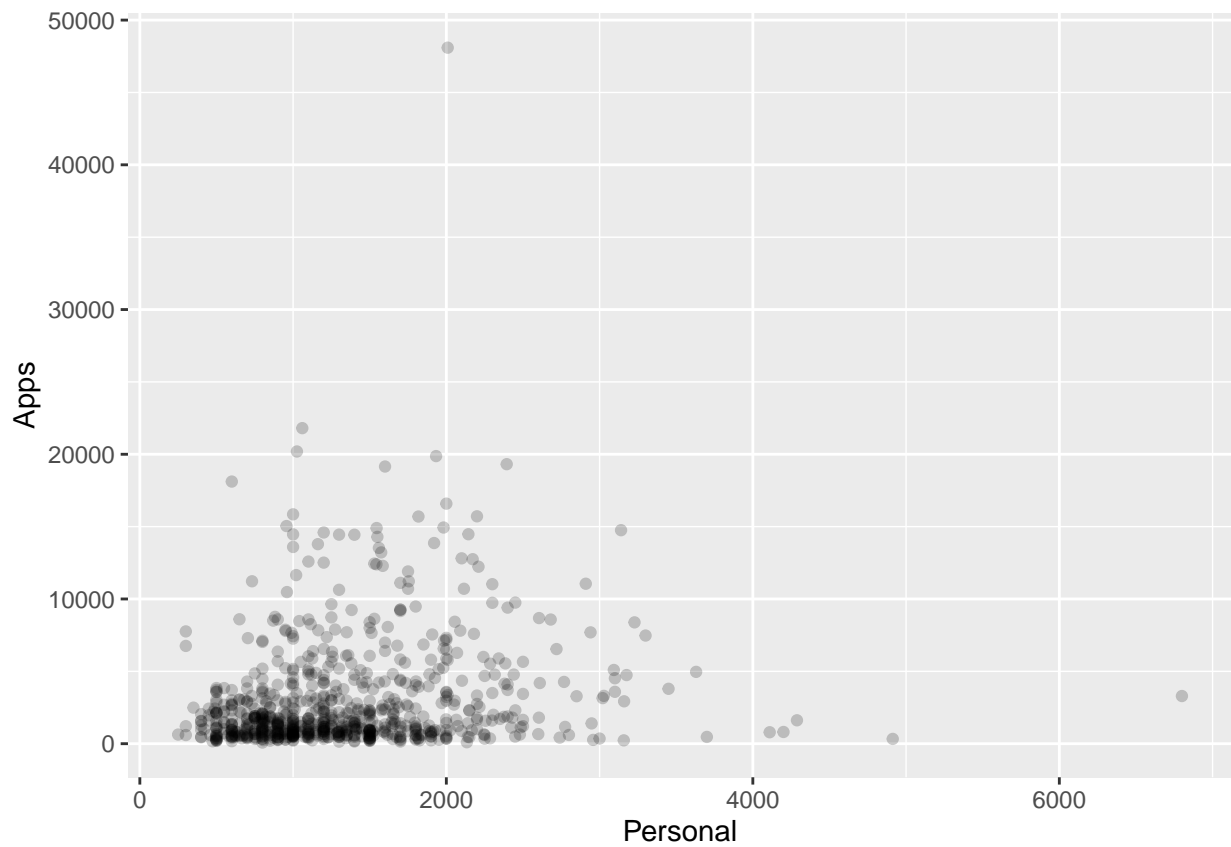


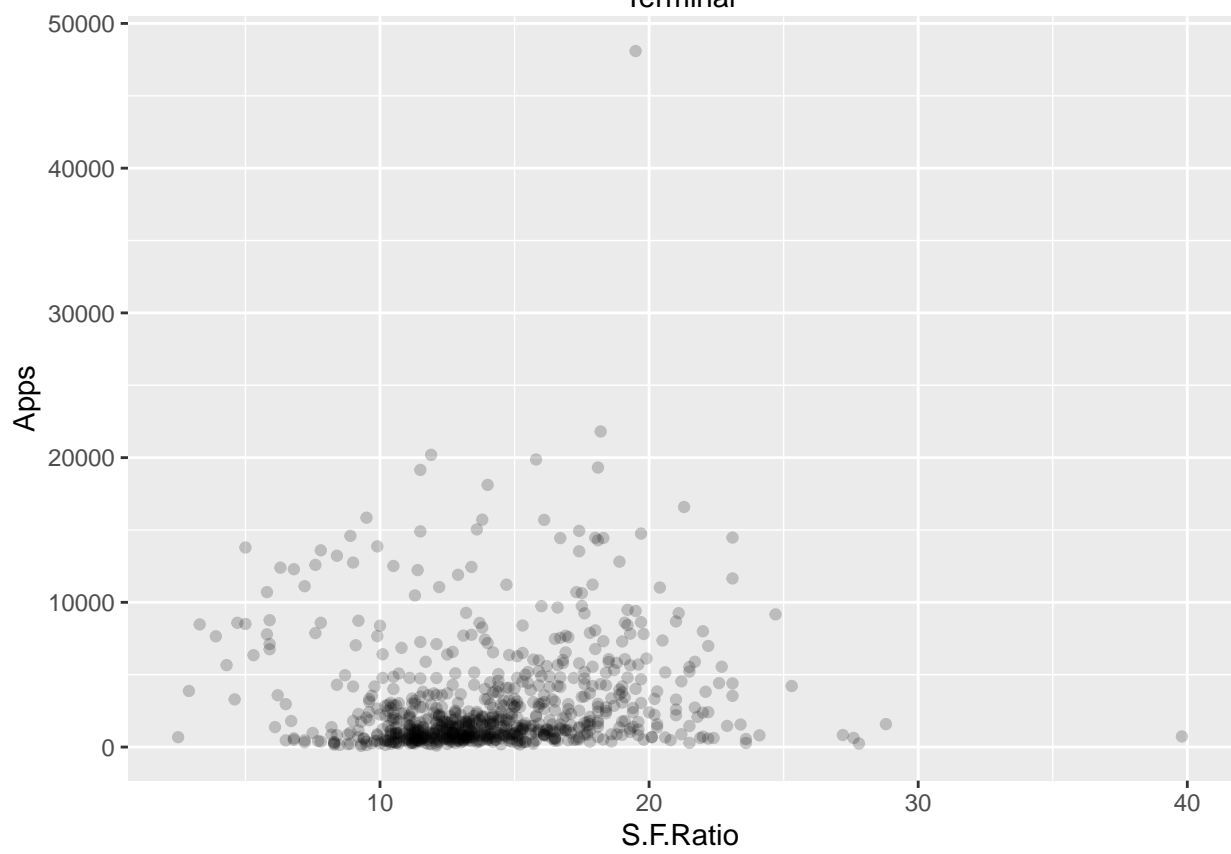
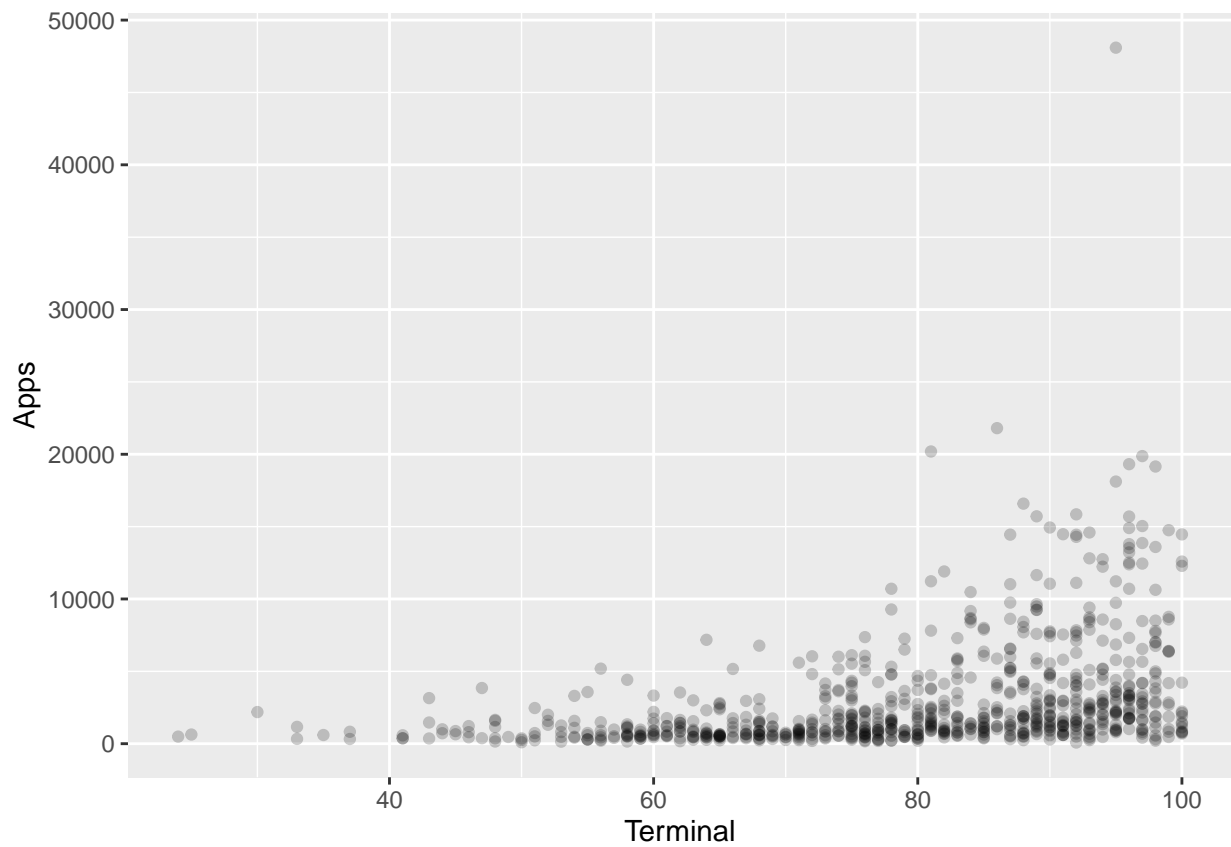


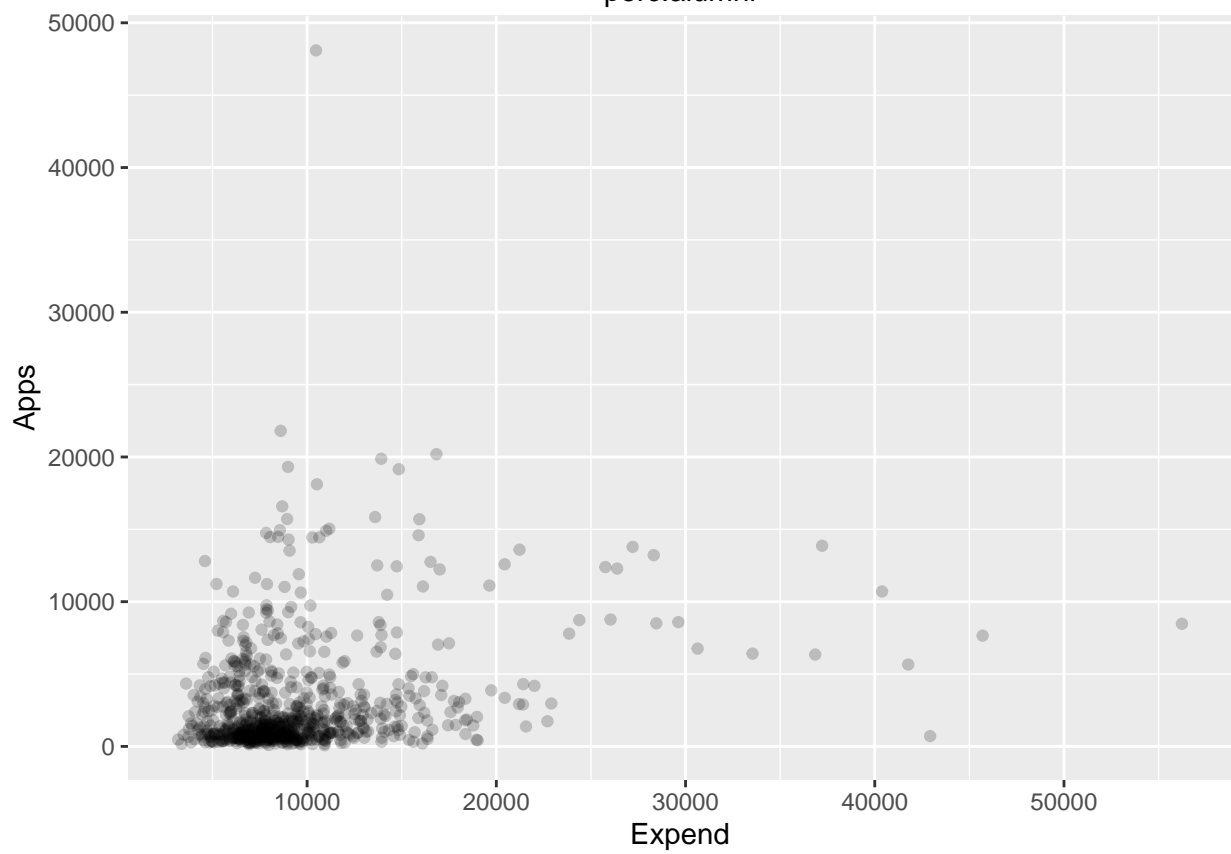
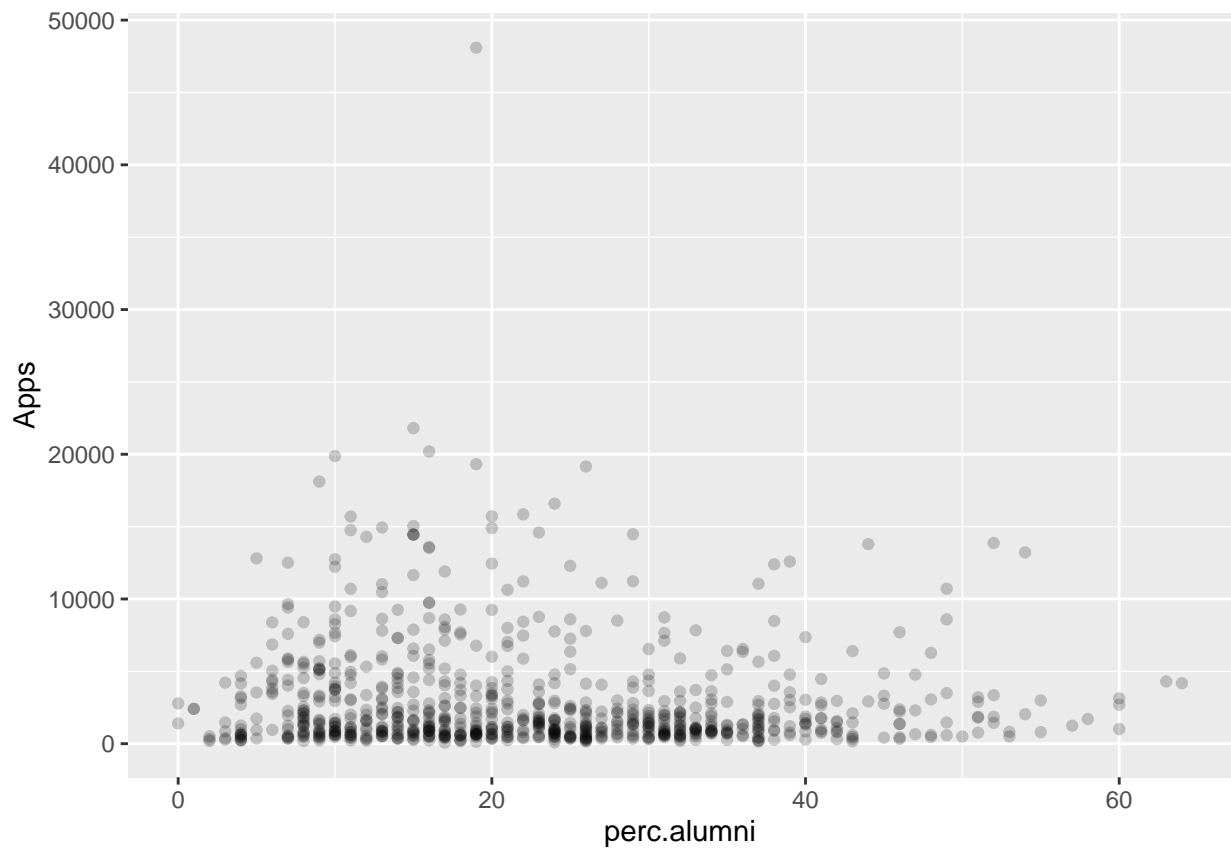


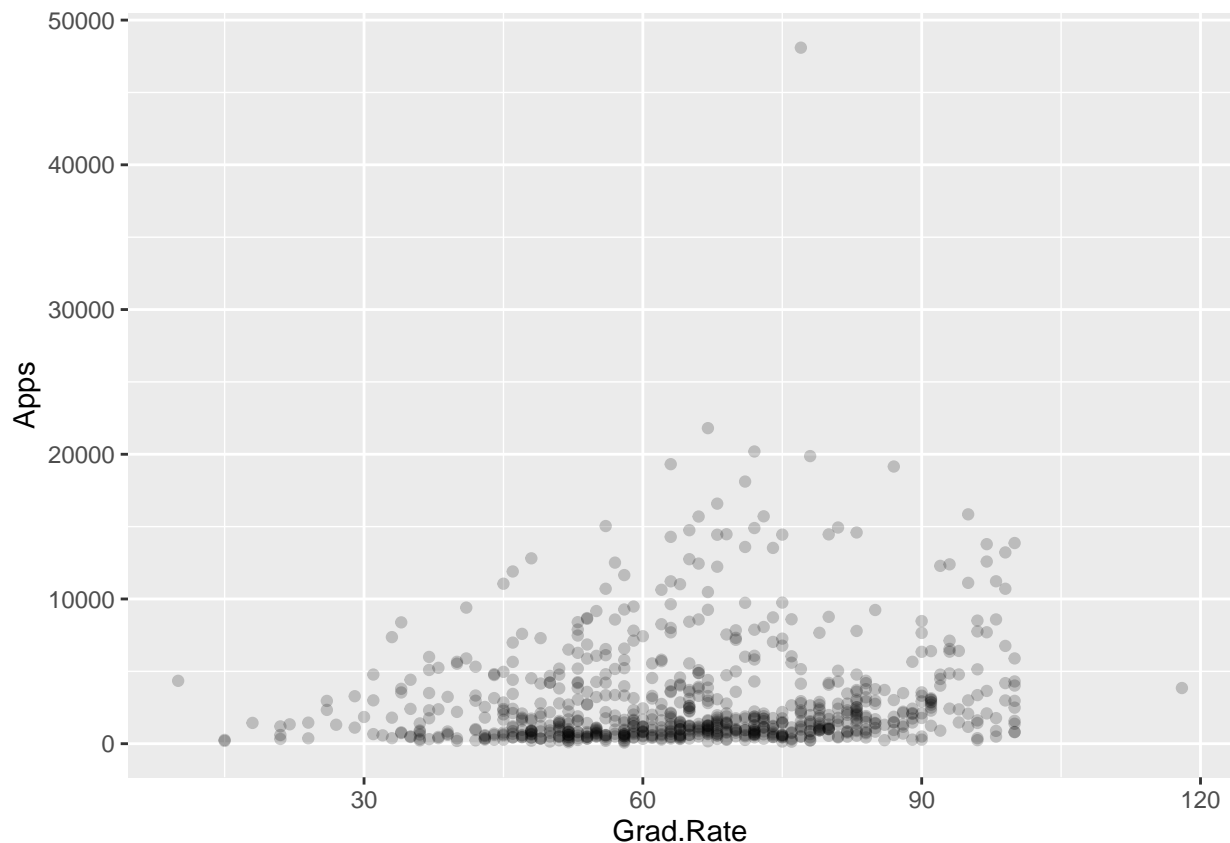












There are some non linear looking relationships. Let's look at books.

```
# Recipe
lasso_pl_recipe <-
  recipe(formula = Apps ~ Books, data = College_train) %>%
  step_poly(all_numeric_predictors(), degree = 2, options = list(raw = TRUE)) %>%
  step_normalize(all_predictors())

# Specification
lasso_pl_spec <-
  linear_reg(mixture = 1, penalty = tune()) %>%
  set_mode("regression") %>%
  set_engine("glmnet")

# Workflow
lasso_pl_workflow <- workflow() %>%
  add_recipe(lasso_pl_recipe) %>%
  add_model(lasso_pl_spec)

# Grid
penalty_grid <- grid_regular(
  penalty(range = c(-5, 2)),
  levels = 50
)

# Tune model
tune_res <- tune_grid(
  lasso_pl_workflow,
```

```

    resamples = College_fold,
    grid = penalty_grid
  )

# Finalize model
lasso_pl_final <- finalize_workflow(lasso_pl_workflow, best_penalty)
lasso_pl_final_fit <- fit(lasso_pl_final, data = College_train)

# Check RMSE
augment(lasso_pl_final_fit, new_data = College_test) %>%
  rmse(truth = Apps, estimate = .pred)

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard      3434.

books_min <- min(College$Books)
books_max <- max(College$Books)
books_range <- tibble(Books = seq(books_min, books_max))
books_range

## # A tibble: 2,245 x 1
##   Books
##   <int>
## 1    96
## 2    97
## 3    98
## 4    99
## 5   100
## 6   101
## 7   102
## 8   103
## 9   104
## 10  105
## # ... with 2,235 more rows

regression_lines <- bind_cols(
  predict(lasso_pl_final_fit, new_data = books_range),
  books_range
)

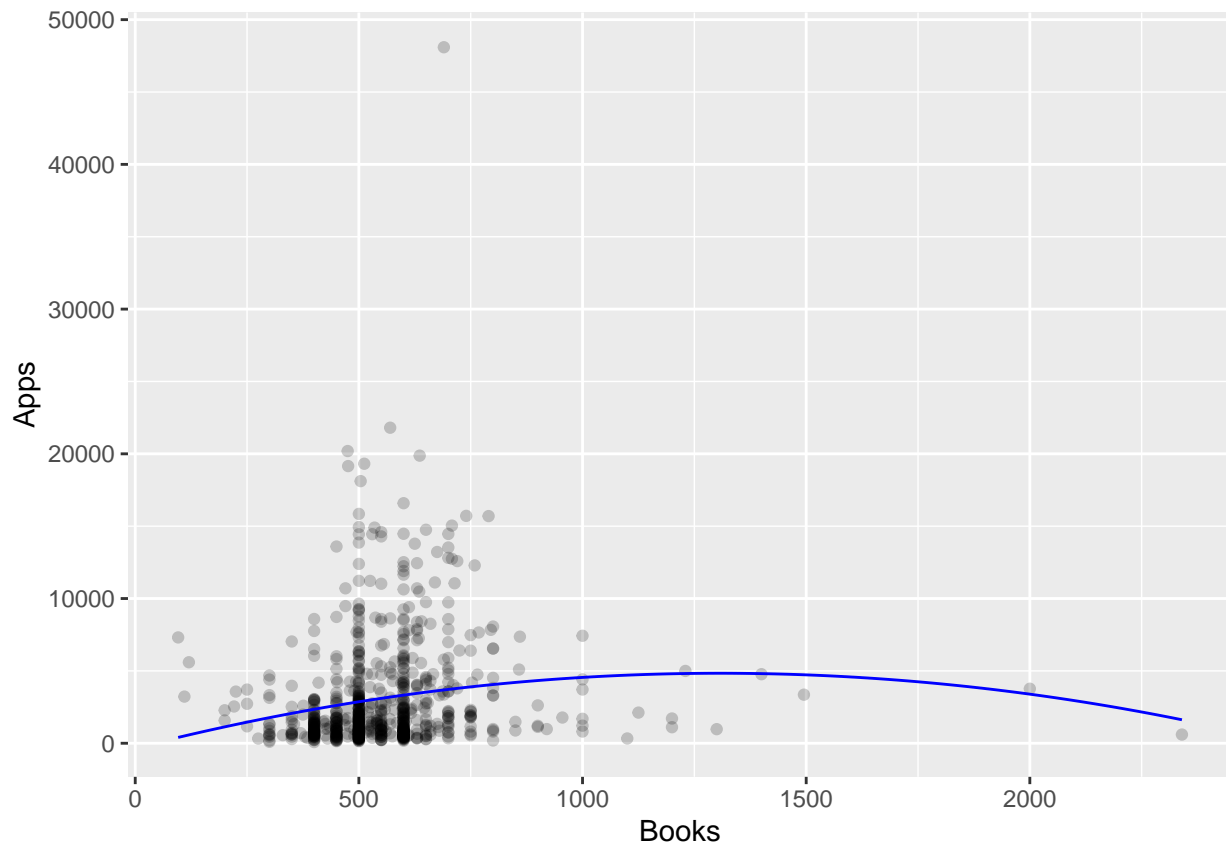
College

## # A tibble: 777 x 18
##   Private Apps Accept Enroll Top10perc Top25perc F.Undergrad P.Undergrad
##   <fct>   <dbl> <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 Yes    1660  1232   721     23     52    2885    537
## 2 Yes    2186  1924   512     16     29    2683   1227
## 3 Yes    1428  1097   336     22     50    1036     99
## 4 Yes     417   349   137     60     89     510     63
## 5 Yes     193   146    55     16     44     249    869
## 6 Yes     587   479   158     38     62     678     41
## 7 Yes     353   340   103     17     45     416    230
## 8 Yes    1899  1720   489     37     68    1594     32

```

```
## 9 Yes      1038    839    227      30      63      973      306
## 10 Yes      582    498    172      21      44      799      78
## # ... with 767 more rows, and 10 more variables: Outstate <dbl>,
## #   Room.Board <dbl>, Books <dbl>, Personal <dbl>, PhD <dbl>, Terminal <dbl>,
## #   S.F.Ratio <dbl>, perc.alumni <dbl>, Expend <dbl>, Grad.Rate <dbl>
```

```
College %>%
  ggplot(aes(Books, Apps)) +
  geom_point(alpha = 0.2) +
  geom_line(aes(y = .pred), data = regression_lines, color = "blue")
```



As we can see a polynomial fits this data better than a linear relationship.

Exercise 3

The Wage data set contains a number of features, such as marital status (marital), job class (jobclass), and others. Explore the relationships between some of these predictors and wage, and use non-linear fitting techniques in order to fit flexible models to the data. Create plots of the results obtained, and write a summary of your findings.

```
Wage <- tibble(Wage)
```

```
basic_eda(Wage)
```

```
## Rows: 3,000
## Columns: 11
## $ year      <int> 2006, 2004, 2003, 2003, 2005, 2008, 2009, 2008, 2006, 2004, ~
## $ age       <int> 18, 24, 45, 43, 50, 54, 44, 30, 41, 52, 45, 34, 35, 39, 54, ~
```

```

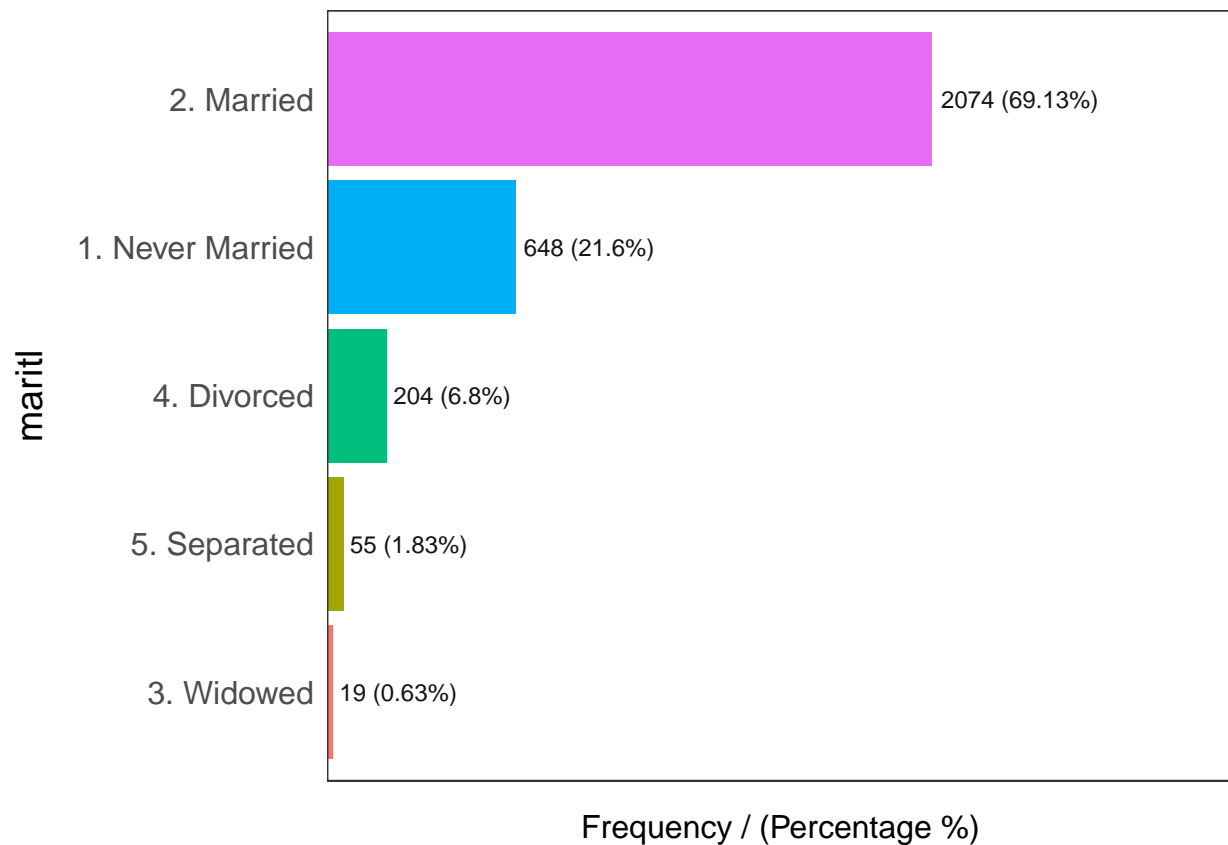
## $ maritl      <fct> 1. Never Married, 1. Never Married, 2. Married, 2. Married,~
## $ race        <fct> 1. White, 1. White, 1. White, 3. Asian, 1. White, 1. White,~
## $ education   <fct> 1. < HS Grad, 4. College Grad, 3. Some College, 4. College ~
## $ region      <fct> 2. Middle Atlantic, 2. Middle Atlantic, 2. Middle Atlantic,~
## $ jobclass     <fct> 1. Industrial, 2. Information, 1. Industrial, 2. Informatio~
## $ health       <fct> 1. <=Good, 2. >=Very Good, 1. <=Good, 2. >=Very Good, 1. <=~
## $ health_ins   <fct> 2. No, 2. No, 1. Yes, 1. Yes, 1. Yes, 1. Yes, 1. Yes, 1. Ye~
## $ logwage      <dbl> 4.318063, 4.255273, 4.875061, 5.041393, 4.318063, 4.845098,~
## $ wage         <dbl> 75.04315, 70.47602, 130.98218, 154.68529, 75.04315, 127.115~
##               variable q_zeros p_zeros q_na p_na q_inf p_inf   type unique
## year           year      0      0    0    0    0    0 integer      7
## age            age      0      0    0    0    0    0 integer     61
## maritl         maritl    0      0    0    0    0    0  factor      5
## race           race      0      0    0    0    0    0  factor      4
## education       education 0      0    0    0    0    0  factor      5
## region         region    0      0    0    0    0    0  factor      1
## jobclass        jobclass 0      0    0    0    0    0  factor      2
## health         health    0      0    0    0    0    0  factor      2
## health_ins     health_ins 0      0    0    0    0    0  factor      2
## logwage        logwage    0      0    0    0    0    0 numeric     508
## wage          wage      0      0    0    0    0    0 numeric     508

## Warning: `guides(<scale> = FALSE)` is deprecated. Please use `guides(<scale> =
## "none")` instead.

##               maritl frequency percentage cumulative_perc
## 1           2. Married      2074      69.13           69.13
## 2 1. Never Married        648      21.60           90.73
## 3      4. Divorced        204       6.80           97.53
## 4      5. Separated         55       1.83           99.36
## 5      3. Widowed         19       0.63          100.00

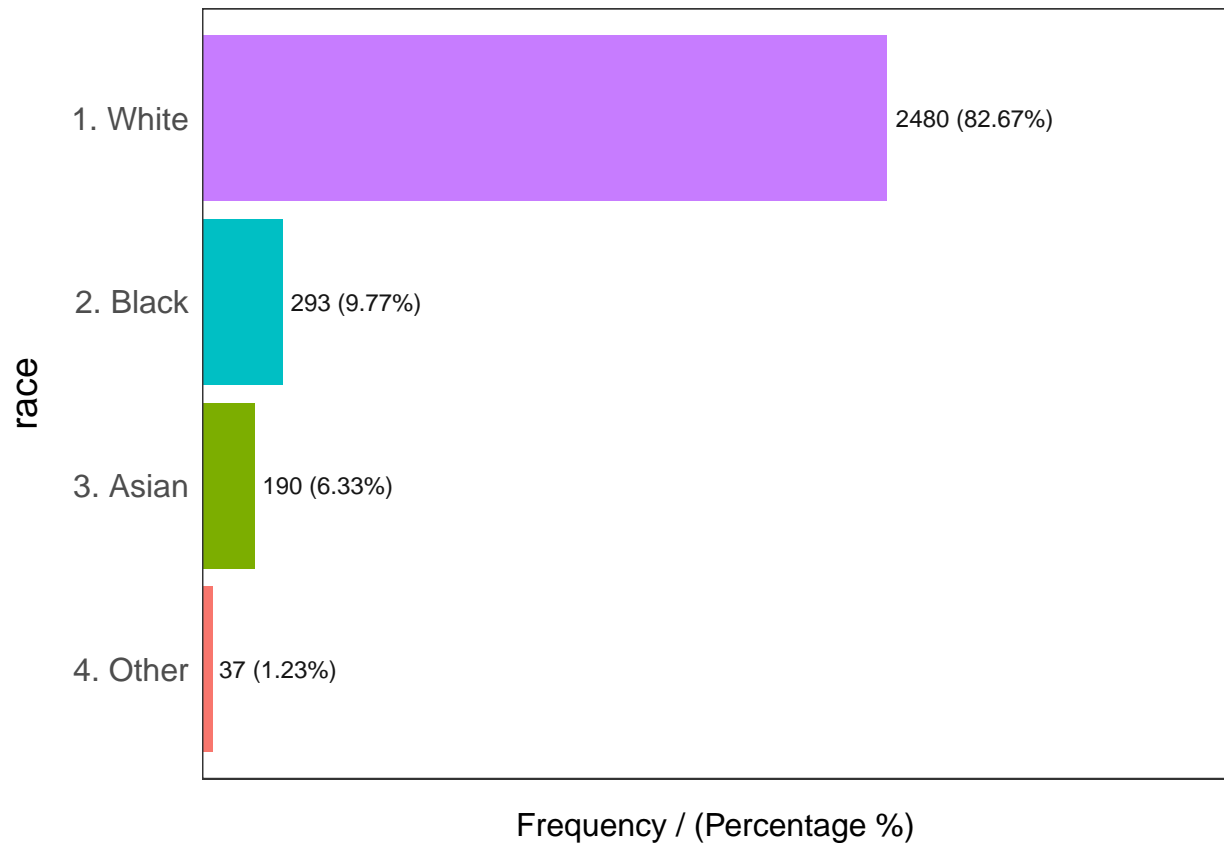
## Warning: `guides(<scale> = FALSE)` is deprecated. Please use `guides(<scale> =
## "none")` instead.

```



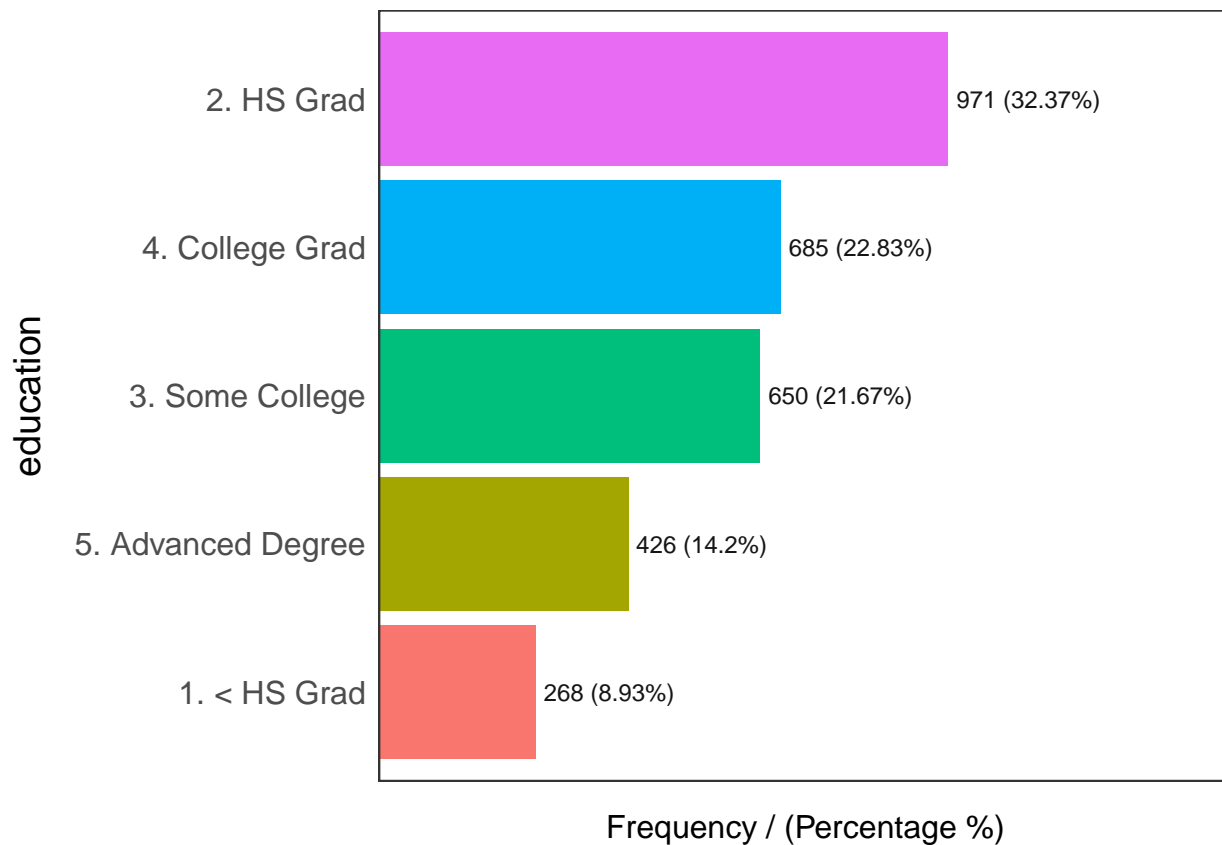
```
##      race frequency percentage cumulative_perc
## 1 1. White      2480      82.67      82.67
## 2 2. Black       293       9.77      92.44
## 3 3. Asian       190       6.33      98.77
## 4 4. Other        37       1.23     100.00

## Warning: `guides(<scale> = FALSE)` is deprecated. Please use `guides(<scale> =
## "none")` instead.
```

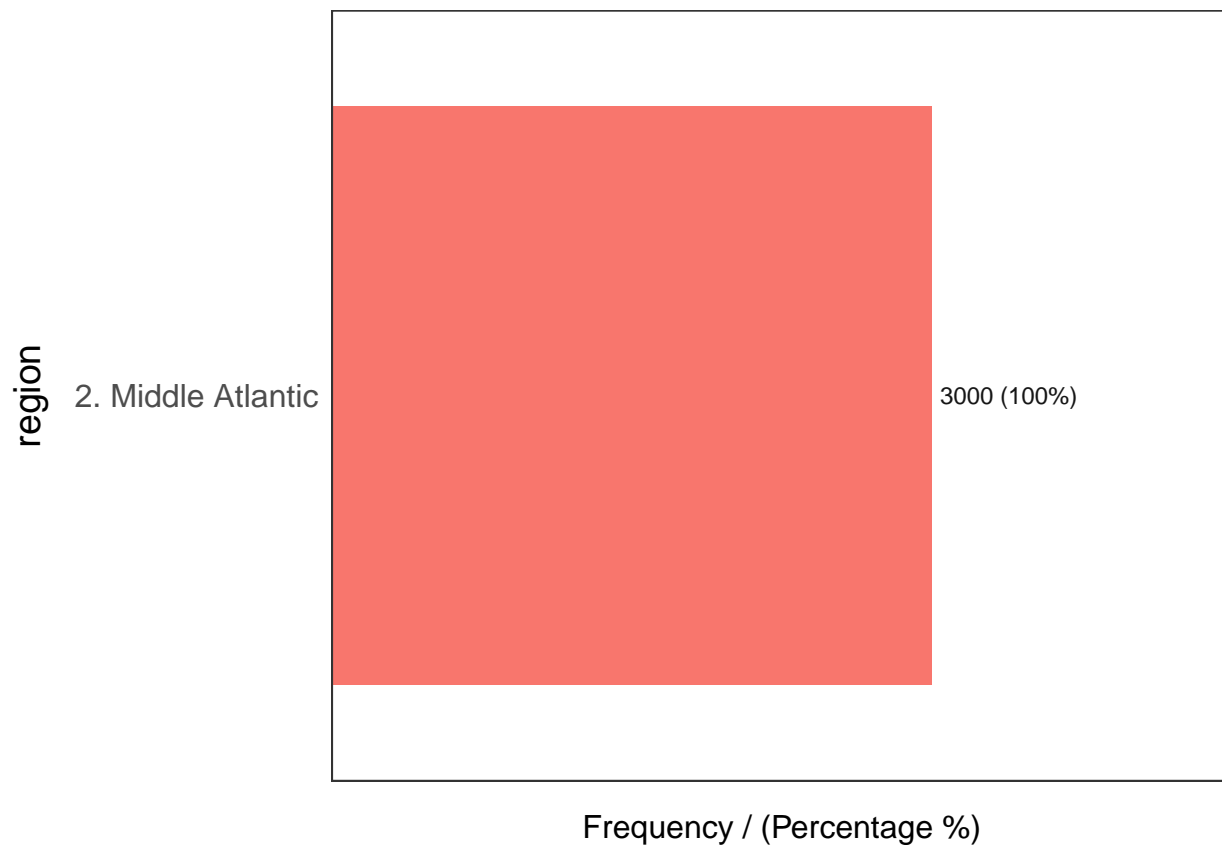


```
##          education frequency percentage cumulative_perc
## 1          2. HS Grad      971         32.37          32.37
## 2          4. College Grad  685         22.83          55.20
## 3          3. Some College  650         21.67          76.87
## 4 5. Advanced Degree      426         14.20          91.07
## 5          1. < HS Grad    268          8.93         100.00

## Warning: `guides(<scale> = FALSE)` is deprecated. Please use `guides(<scale> =
## "none")` instead.
```

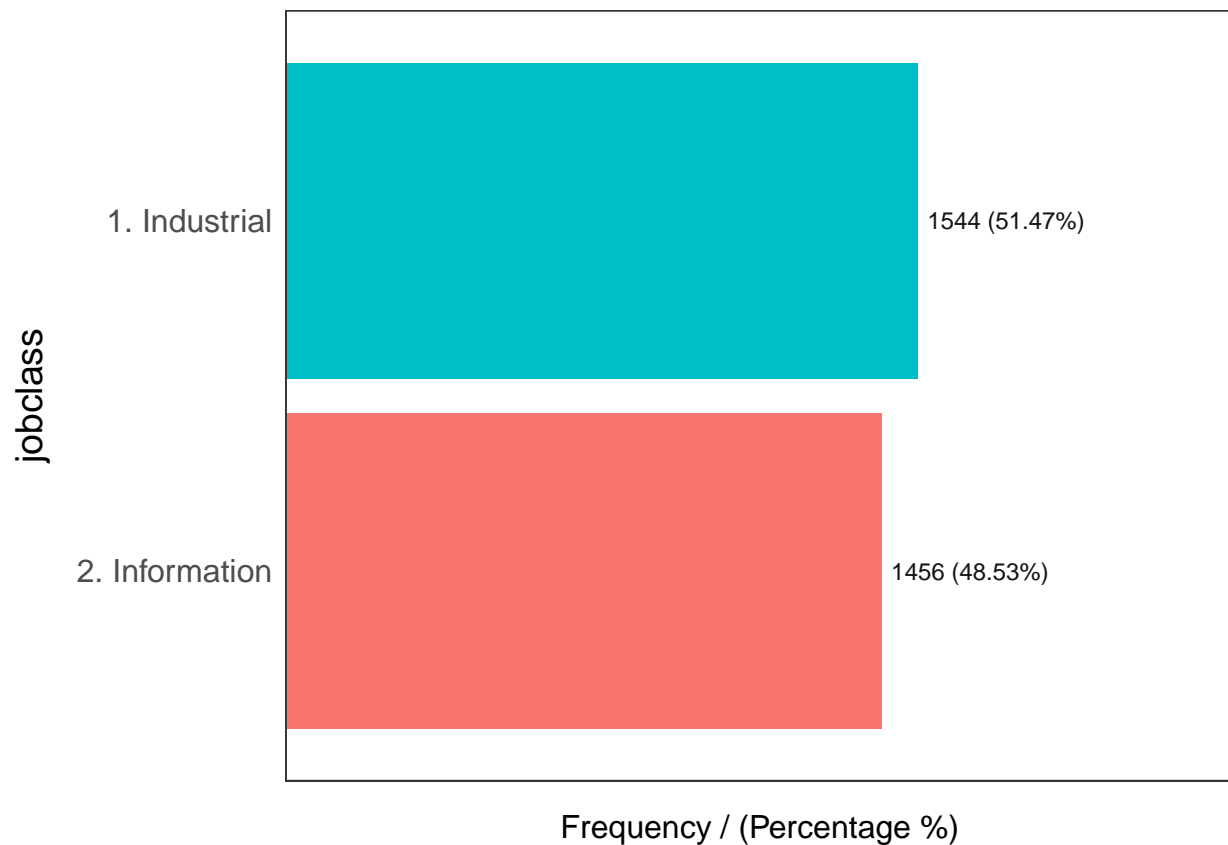



```
##           region frequency percentage cumulative_perc
## 1 2. Middle Atlantic      3000         100         100
## Warning: `guides(<scale> = FALSE)` is deprecated. Please use `guides(<scale> =
## "none")` instead.
```



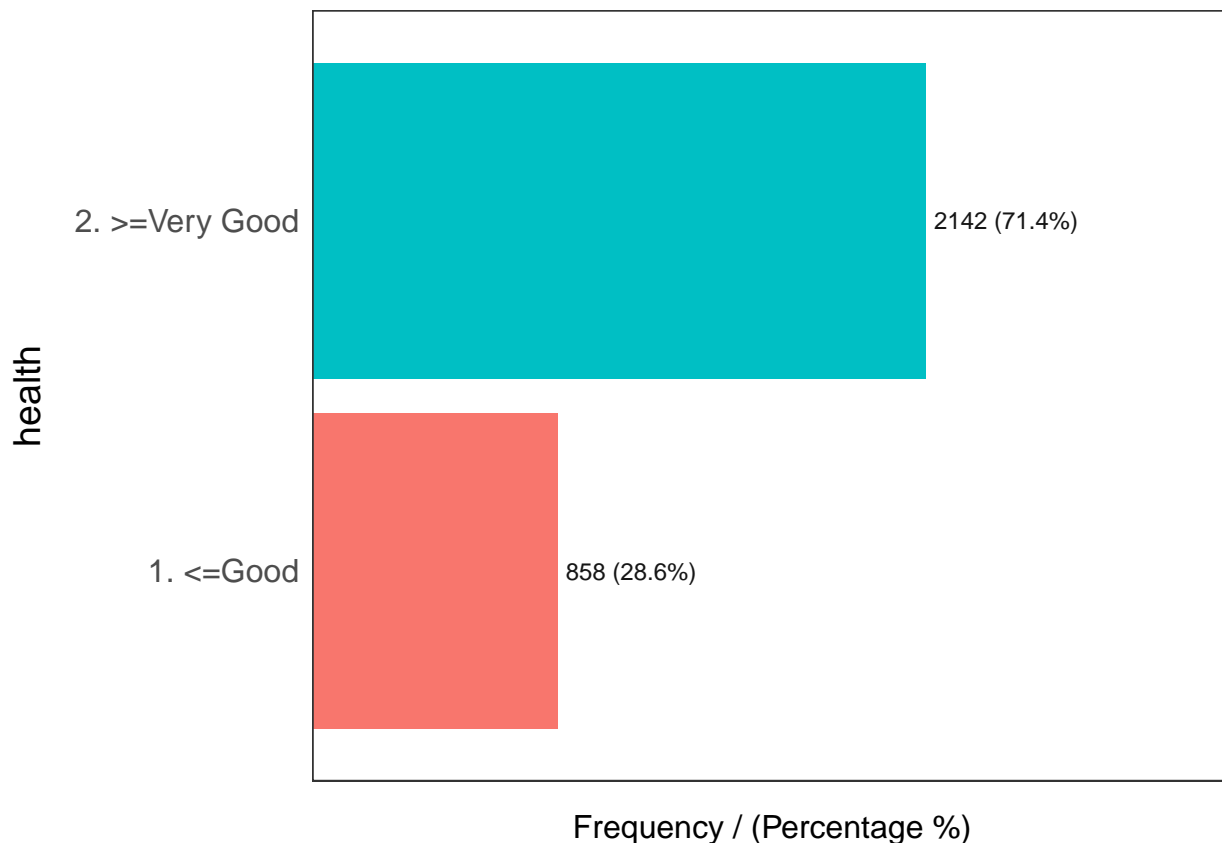
```
##           jobclass frequency percentage cumulative_perc
## 1  1. Industrial      1544      51.47          51.47
## 2  2. Information      1456      48.53          100.00

## Warning: `guides(<scale> = FALSE)` is deprecated. Please use `guides(<scale> =
## "none")` instead.
```



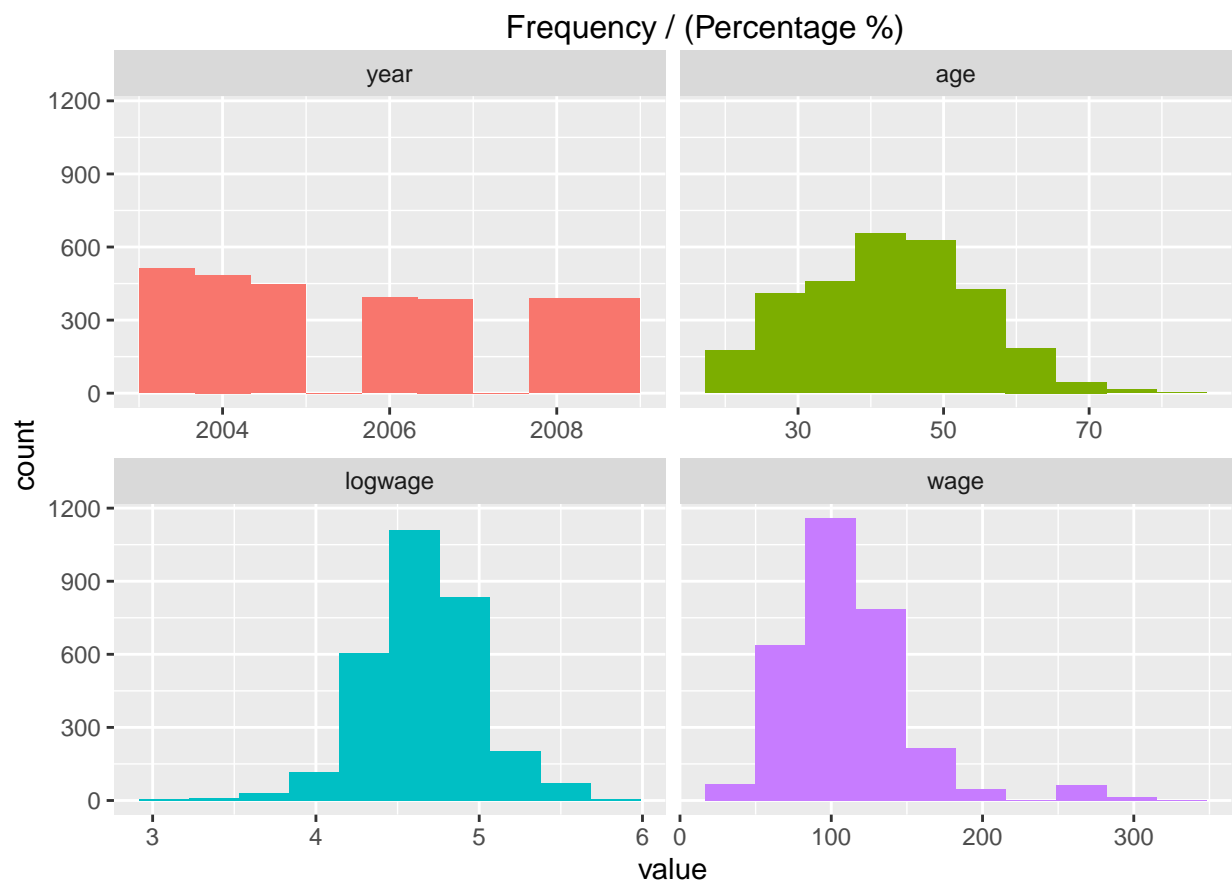
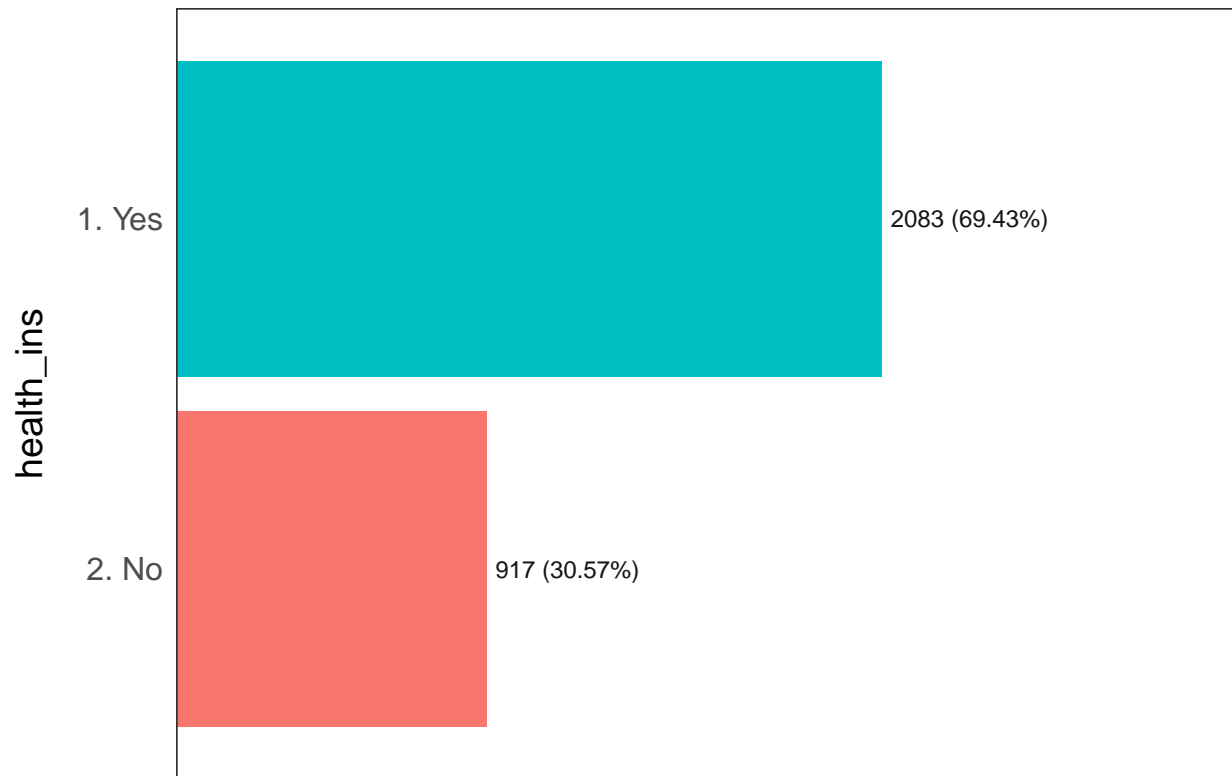
```
##           health frequency percentage cumulative_perc
## 1 2. >=Very Good      2142         71.4           71.4
## 2   1. <=Good         858         28.6           100.0

## Warning: `guides(<scale> = FALSE)` is deprecated. Please use `guides(<scale> =
## "none")` instead.
```



```
## health_ins frequency percentage cumulative_perc
## 1 1. Yes 2083 69.43 69.43
## 2 2. No 917 30.57 100.00
##
## variable mean std_dev variation_coef p_01 p_05
## 1 year 2005.791000 2.0261673 0.001010159 2003.000000 2003.000000
## 2 age 42.414667 11.5424056 0.272132414 20.000000 24.000000
## 3 logwage 4.653905 0.3517526 0.075582244 3.698918 4.113943
## 4 wage 111.703608 41.7285955 0.373565332 40.403552 61.187526
## p_25 p_50 p_75 p_95 p_99 skewness
## 1 2004.000000 2006.000000 2008.000000 2009.000000 2009.000000 0.1428111
## 2 33.750000 42.000000 51.000000 61.000000 70.000000 0.1477340
## 3 4.447158 4.653213 4.857332 5.176091 5.626186 -0.1235535
## 4 85.383940 104.921507 128.680488 176.989650 277.601418 1.6814889
## kurtosis iqr range_98
## 1 1.733853 4.0000000 [2003, 2009]
## 2 2.552129 17.2500000 [20, 70]
## 3 4.728038 0.4101745 [3.6989175737819, 5.62618633492728]
## 4 7.828952 43.2965478 [40.4035523122557, 277.601417511009]
## range_80
## 1 [2003, 2009]
## 2 [27, 58]
## 3 [4.25527250510331, 5.04151096788643]
## 4 [70.4760196469445, 154.703600419223]

## Warning: `guides(<scale> = FALSE)` is deprecated. Please use `guides(<scale> =
## "none")` instead.
```



data

```

##
## 11 Variables      3000 Observations
## -----
## year
##      n missing distinct      Info      Mean      Gmd
##    3000      0      7    0.979    2006    2.312
##
## lowest : 2003 2004 2005 2006 2007, highest: 2005 2006 2007 2008 2009
##
## Value      2003  2004  2005  2006  2007  2008  2009
## Frequency   513   485   447   392   386   388   389
## Proportion 0.171 0.162 0.149 0.131 0.129 0.129 0.130
## -----
## age
##      n missing distinct      Info      Mean      Gmd      .05      .10
##    3000      0      61    0.999    42.41    13.16    24.00    27.00
##      .25      .50      .75      .90      .95
##    33.75    42.00    51.00    58.00    61.00
##
## lowest : 18 19 20 21 22, highest: 74 75 76 77 80
## -----
## maritl
##      n missing distinct
##    3000      0      5
##
## lowest : 1. Never Married 2. Married      3. Widowed      4. Divorced      5. Separated
## highest: 1. Never Married 2. Married      3. Widowed      4. Divorced      5. Separated
##
## Value      1. Never Married      2. Married      3. Widowed      4. Divorced
## Frequency           648           2074           19           204
## Proportion           0.216           0.691           0.006           0.068
##
## Value      5. Separated
## Frequency           55
## Proportion           0.018
## -----
## race
##      n missing distinct
##    3000      0      4
##
## Value      1. White 2. Black 3. Asian 4. Other
## Frequency      2480      293      190      37
## Proportion      0.827      0.098      0.063      0.012
## -----
## education
##      n missing distinct
##    3000      0      5
##
## lowest : 1. < HS Grad      2. HS Grad      3. Some College      4. College Grad      5. Advanced Deg
## highest: 1. < HS Grad      2. HS Grad      3. Some College      4. College Grad      5. Advanced Deg
##
## Value      1. < HS Grad      2. HS Grad      3. Some College
## Frequency           268           971           650
## Proportion           0.089           0.324           0.217

```

```

##
## Value          4. College Grad 5. Advanced Degree
## Frequency          685          426
## Proportion          0.228          0.142
## -----
## region
##          n          missing          distinct          value
##          3000          0          1 2. Middle Atlantic
##
## Value          2. Middle Atlantic
## Frequency          3000
## Proportion          1
## -----
## jobclass
##          n missing distinct
##          3000          0          2
##
## Value          1. Industrial 2. Information
## Frequency          1544          1456
## Proportion          0.515          0.485
## -----
## health
##          n missing distinct
##          3000          0          2
##
## Value          1. <=Good 2. >=Very Good
## Frequency          858          2142
## Proportion          0.286          0.714
## -----
## health_ins
##          n missing distinct
##          3000          0          2
##
## Value          1. Yes 2. No
## Frequency          2083          917
## Proportion          0.694          0.306
## -----
## logwage
##          n missing distinct          Info          Mean          Gmd          .05          .10
##          3000          0          508          1          4.654          0.3824          4.114          4.255
##          .25          .50          .75          .90          .95
##          4.447          4.653          4.857          5.042          5.176
##
## lowest : 3.000000 3.041393 3.133858 3.147367 3.176091
## highest: 5.701323 5.735190 5.742793 5.750441 5.763128
## -----
## wage
##          n missing distinct          Info          Mean          Gmd          .05          .10
##          3000          0          508          1          111.7          42.64          61.19          70.48
##          .25          .50          .75          .90          .95
##          85.38          104.92          128.68          154.70          176.99
##
## lowest : 20.08554 20.93438 22.96240 23.27470 23.95294
## highest: 299.26298 309.57177 311.93457 314.32934 318.34243

```

```
## -----
Wage_split <- initial_split(Wage, strata = wage, prop = 0.5)
Wage_split

## <Analysis/Assess/Total>
## <1499/1501/3000>

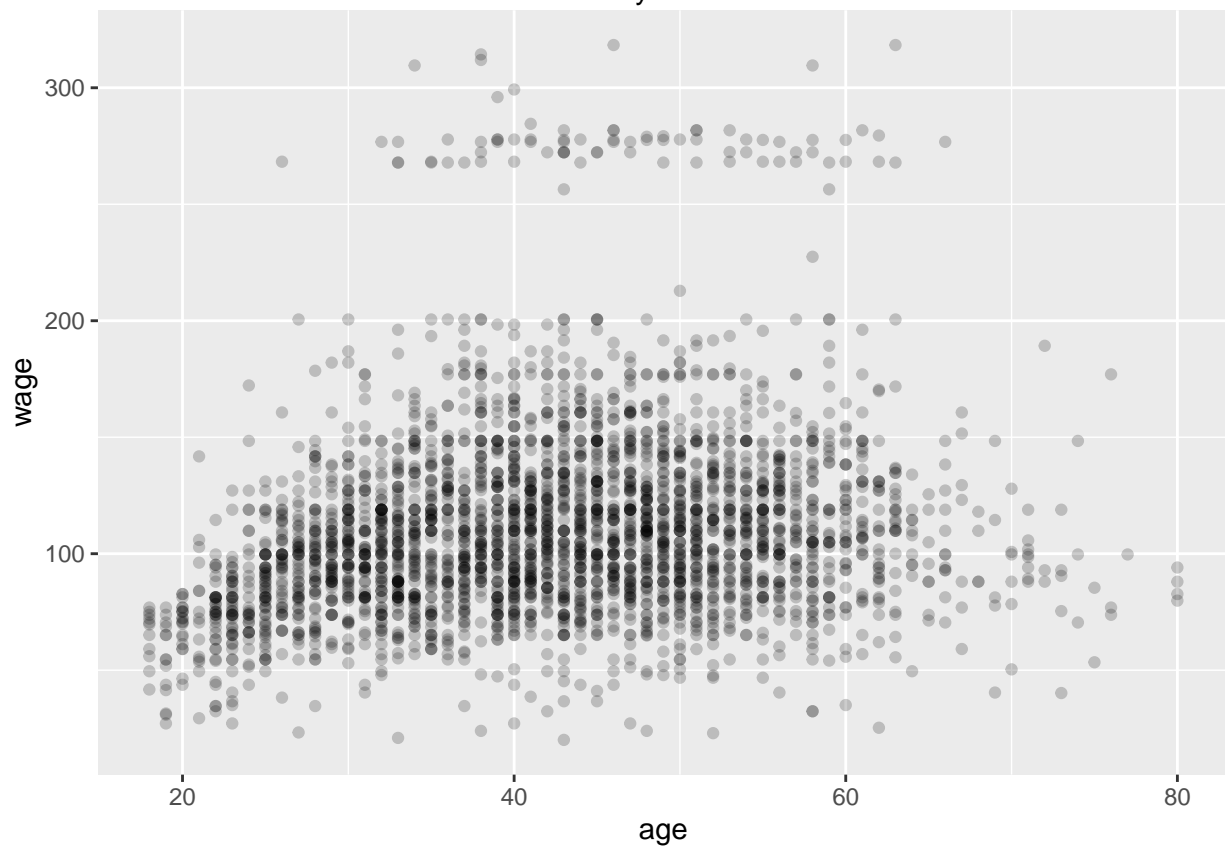
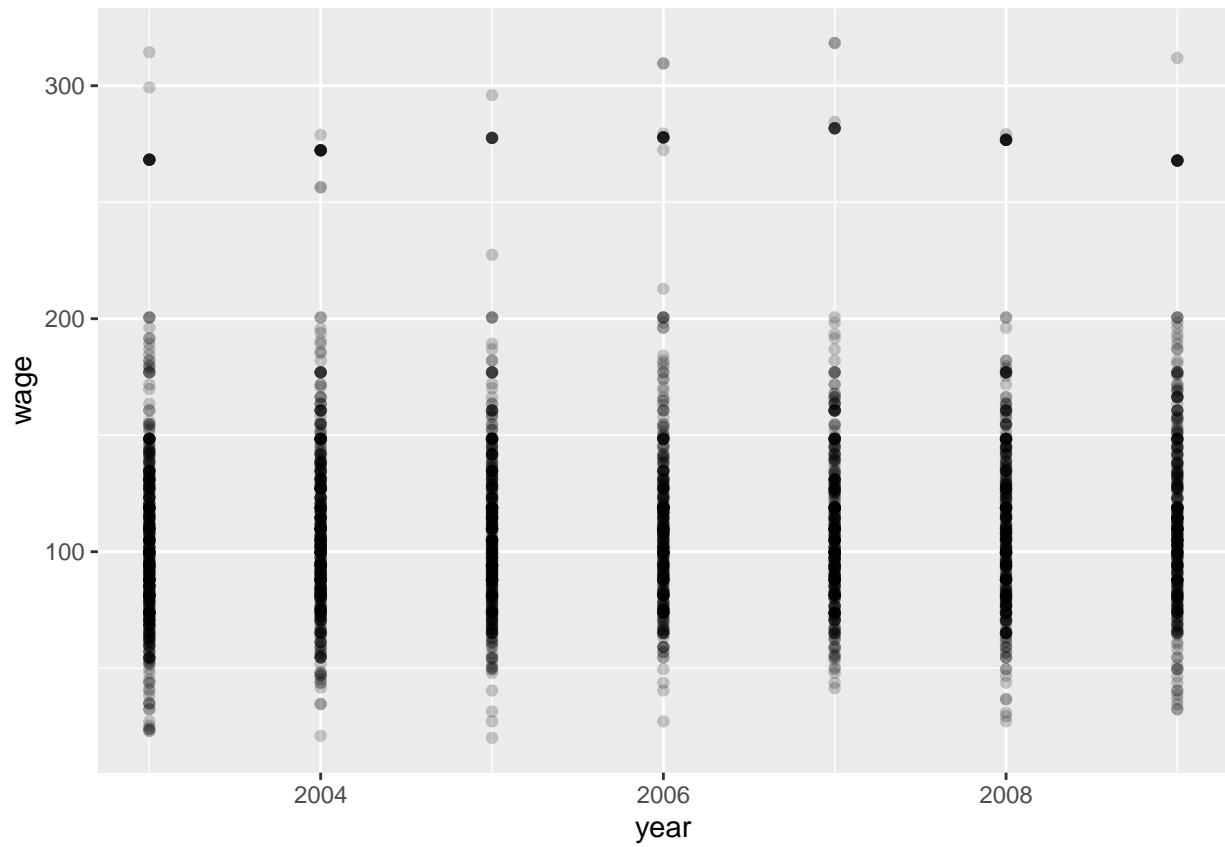
Wage_train <- training(Wage_split)
Wage_test <- testing(Wage_split)

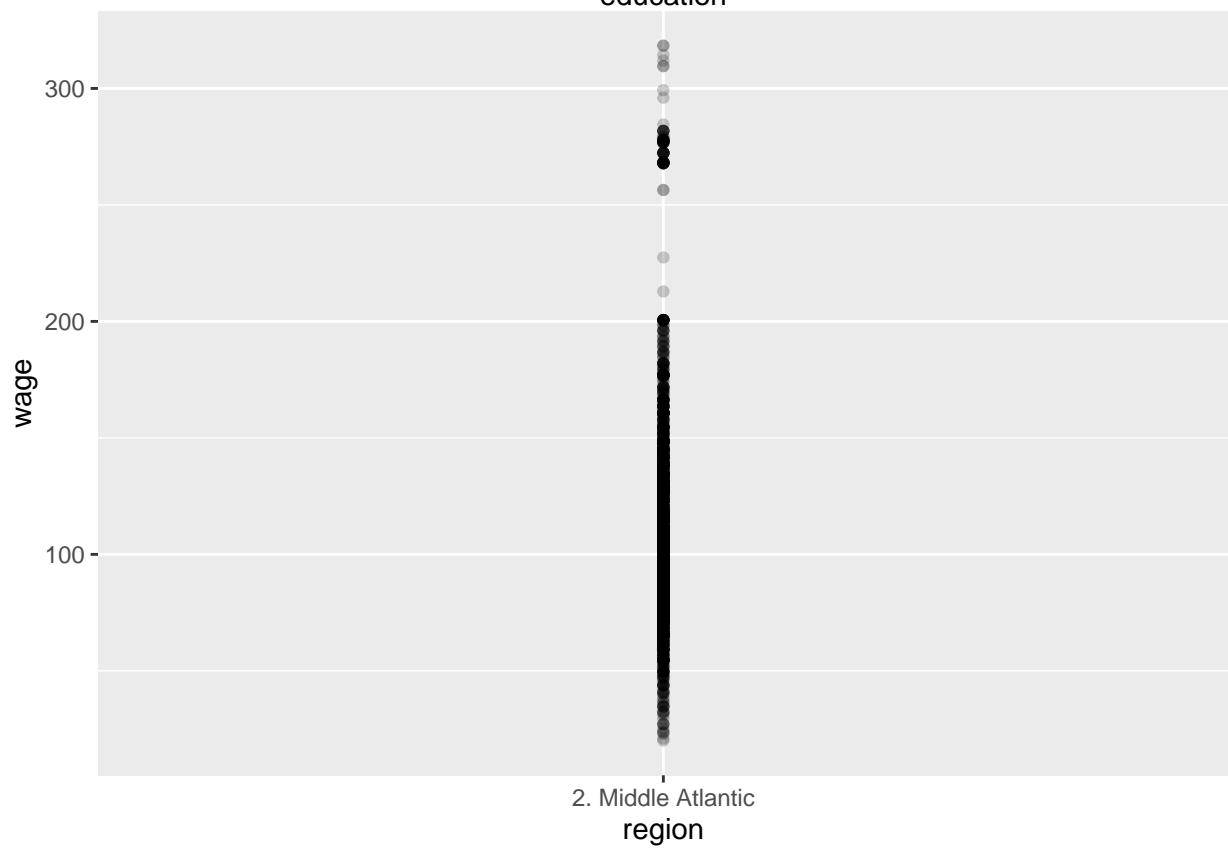
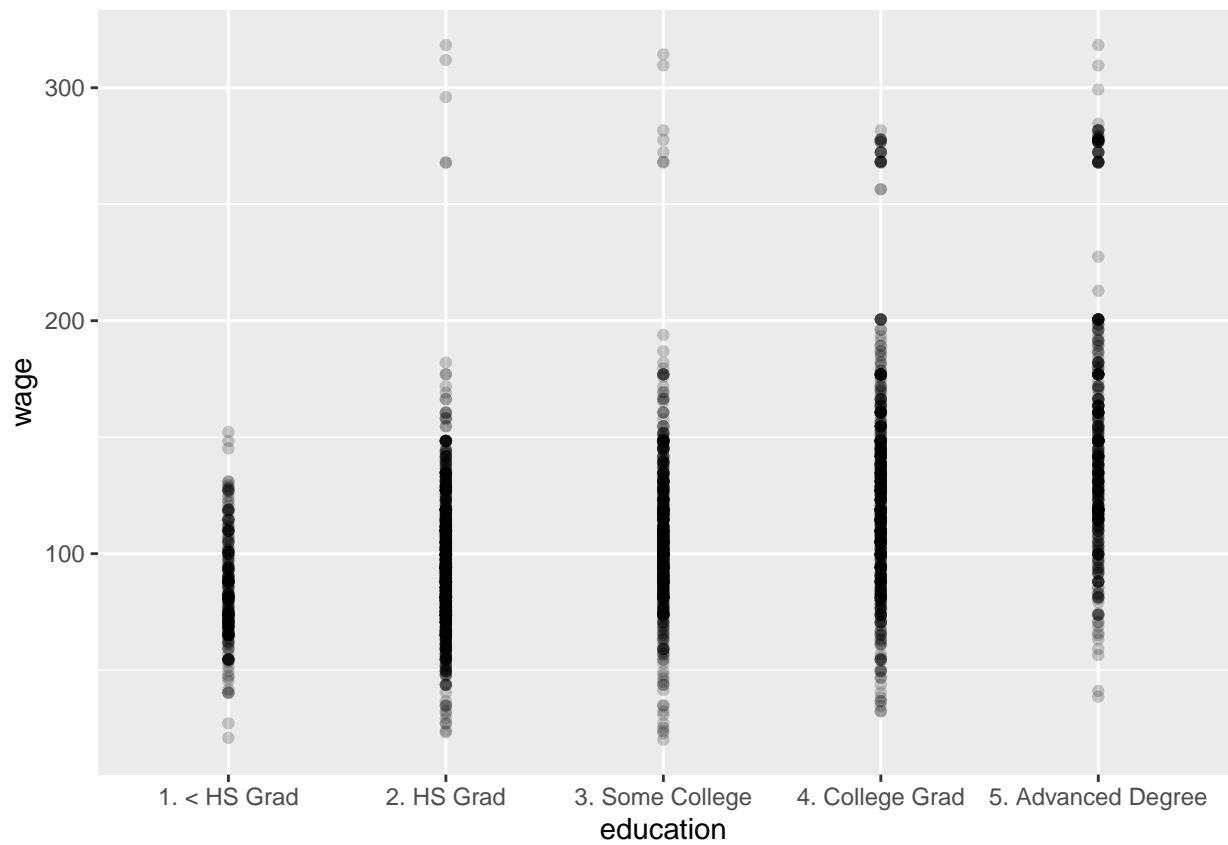
Wage_fold <- vfold_cv(Wage_train, v = 10)
Wage_fold

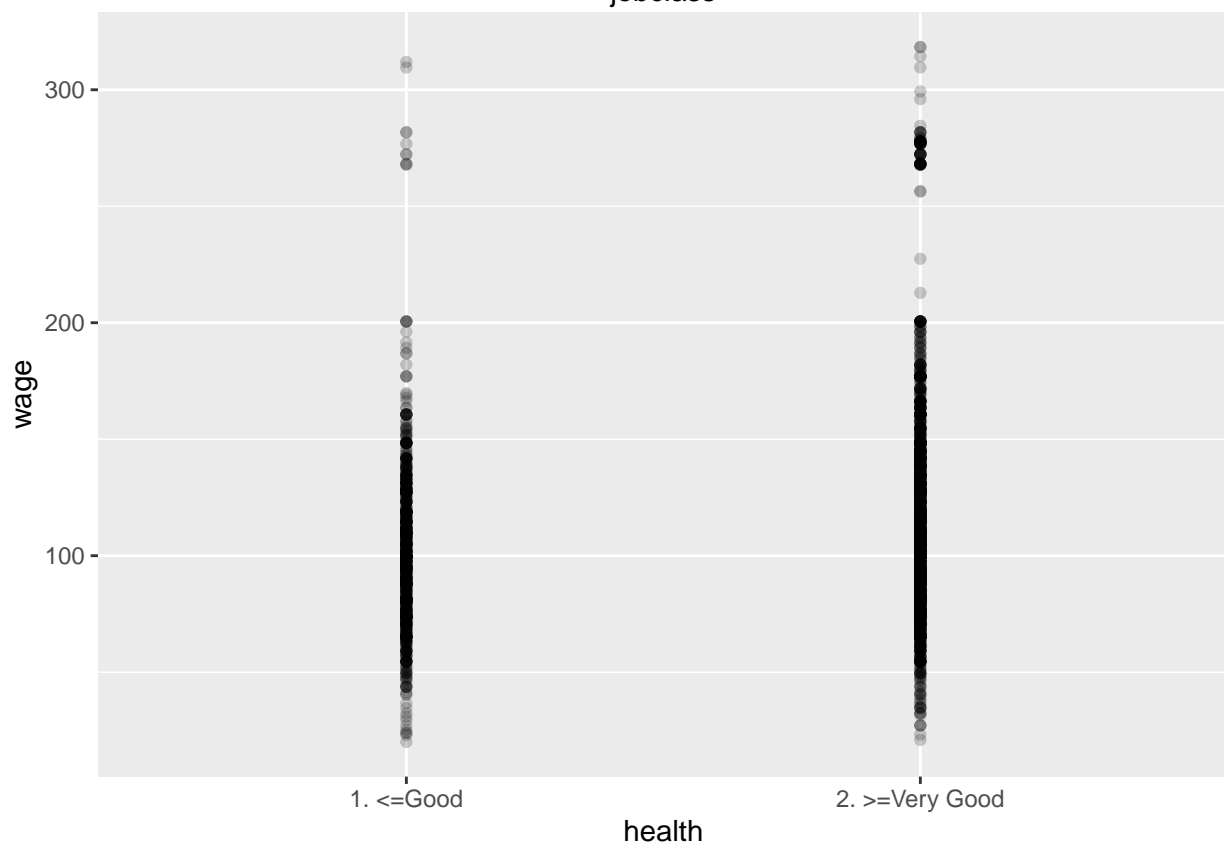
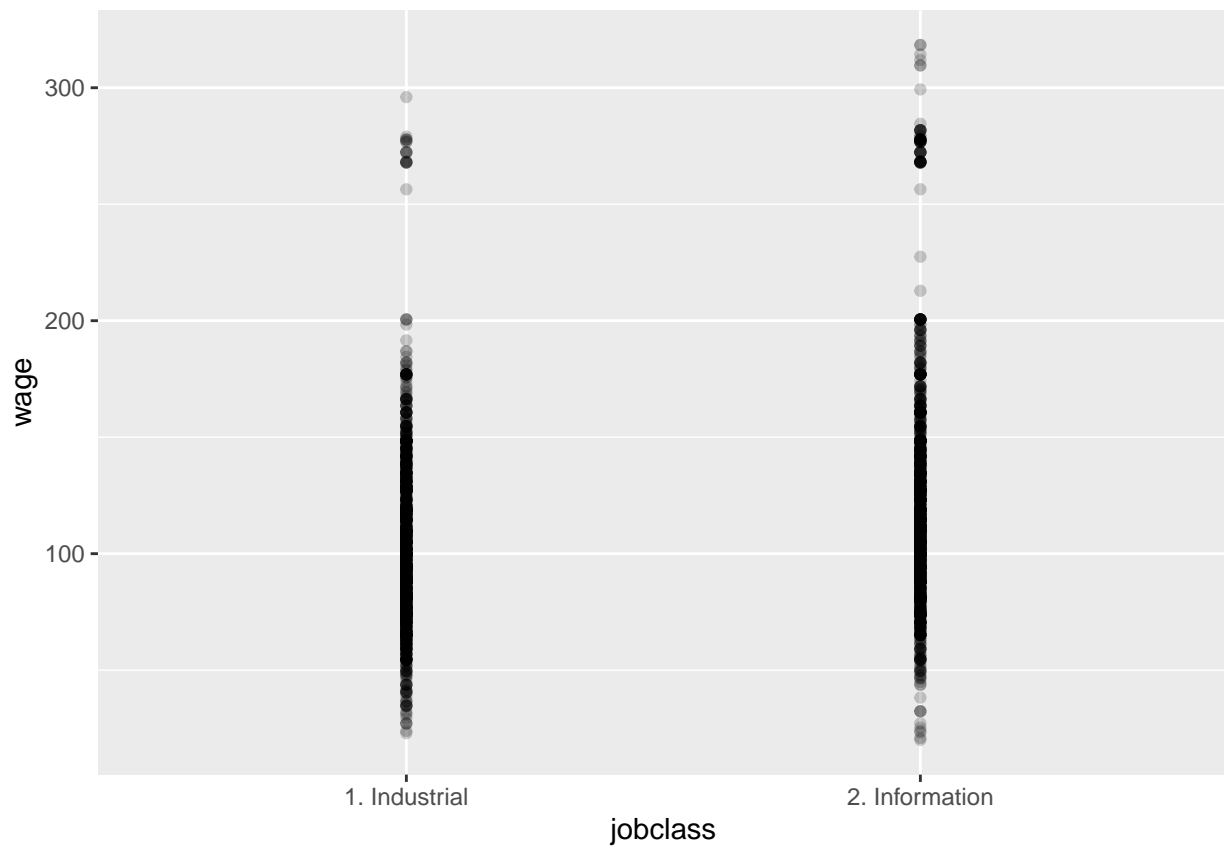
## # 10-fold cross-validation
## # A tibble: 10 x 2
##   splits      id
##   <list>      <chr>
## 1 <split [1349/150]> Fold01
## 2 <split [1349/150]> Fold02
## 3 <split [1349/150]> Fold03
## 4 <split [1349/150]> Fold04
## 5 <split [1349/150]> Fold05
## 6 <split [1349/150]> Fold06
## 7 <split [1349/150]> Fold07
## 8 <split [1349/150]> Fold08
## 9 <split [1349/150]> Fold09
## 10 <split [1350/149]> Fold10

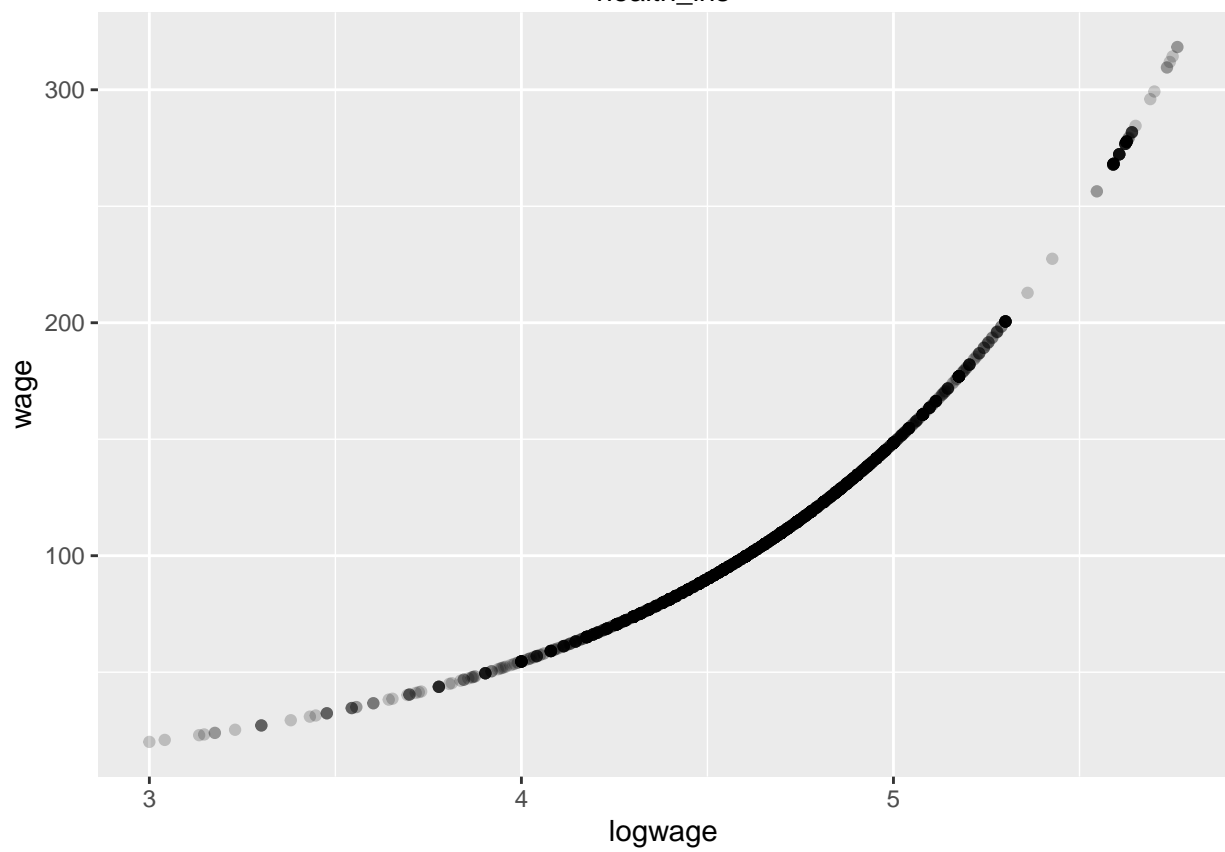
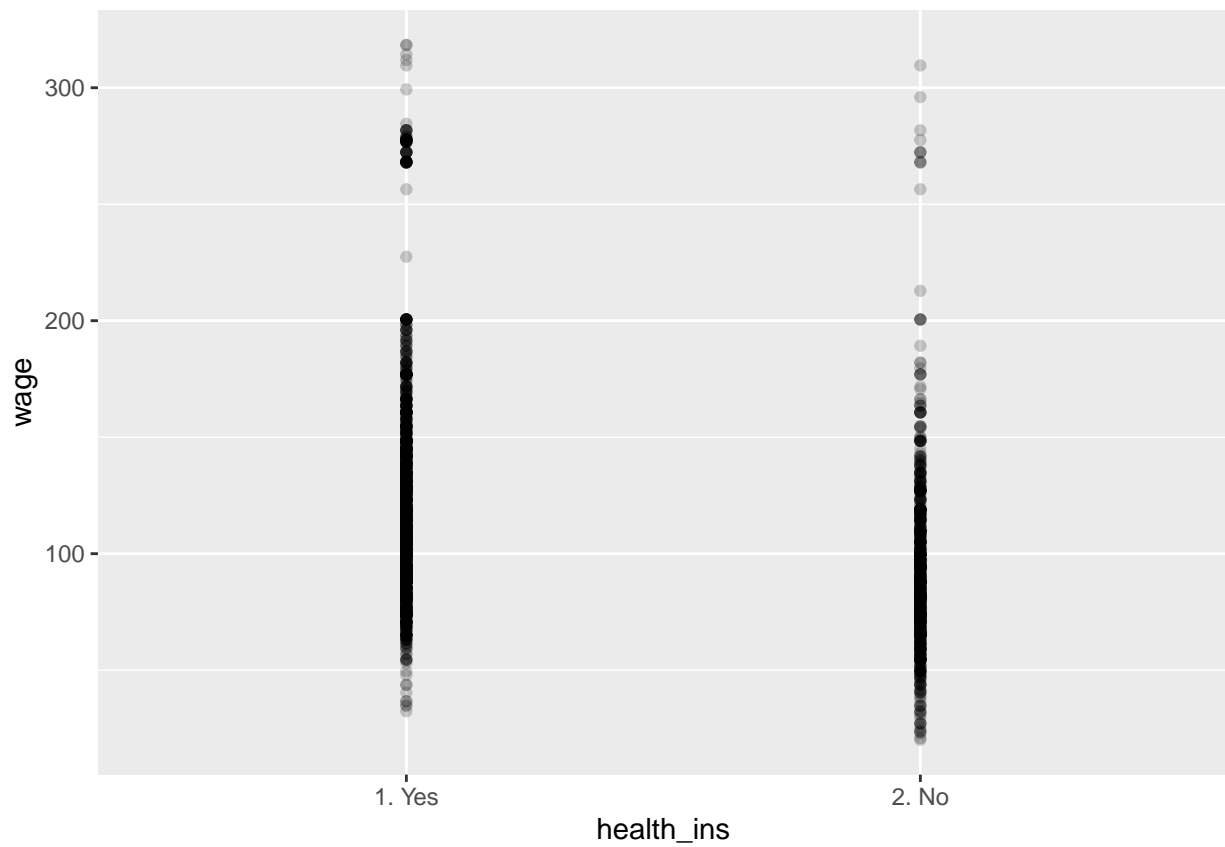
colnames <- colnames(Wage)
colnames <- colnames[colnames != "wage"]

for (colname in colnames) {
  pl <- Wage %>%
    ggplot(aes_string(x = colname, y = "wage")) +
    geom_point(alpha = 0.2)
  print(pl)
}
```







```

# Recipe
lm_pl_recipe <-
  recipe(formula = wage ~ age, data = Wage_train) %>%
  step_poly(age, degree = 2, options = list(raw = TRUE)) %>%
  step_normalize(all_predictors())

# Specification
lm_pl_spec <-
  linear_reg() %>%
  set_mode("regression") %>%
  set_engine("lm")

# Workflow
lm_pl_workflow <- workflow() %>%
  add_recipe(lm_pl_recipe) %>%
  add_model(lm_pl_spec)

# Finalize model
lm_pl_fit <- fit(lm_pl_workflow, data = Wage_train)

# Check RMSE
augment(lm_pl_fit, new_data = Wage_test) %>%
  rmse(truth = wage, estimate = .pred)

```

```

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard      39.8

```

```
tidy(lm_pl_fit)
```

```

## # A tibble: 3 x 5
##   term          estimate std.error statistic  p.value
##   <chr>         <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)   112.      1.04    108.      0
## 2 age_poly_1    62.9      6.67     9.43 1.52e-20
## 3 age_poly_2   -56.1      6.67    -8.40 1.01e-16

```

```

age_min <- min(Wage$age)
age_max <- max(Wage$age)
age_range <- tibble(age = seq(age_min, age_max))
age_range

```

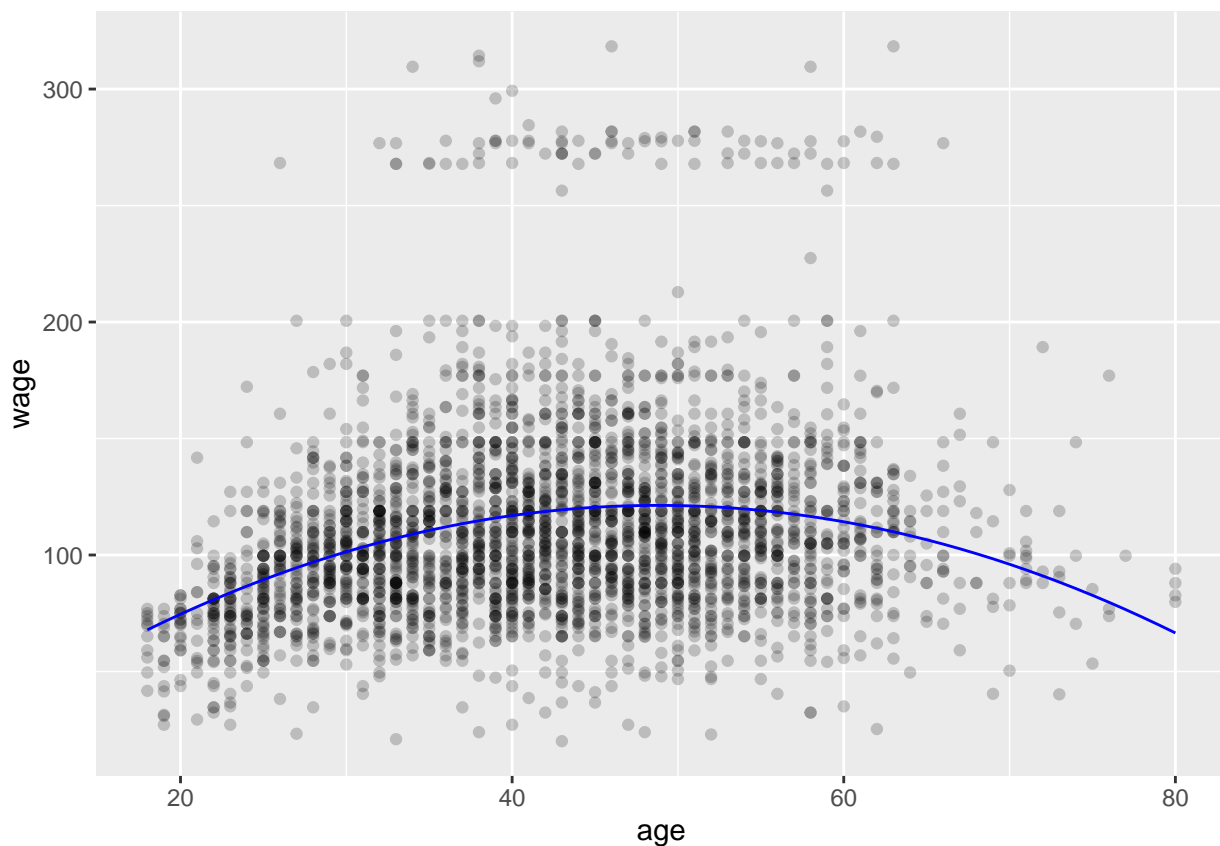
```

## # A tibble: 63 x 1
##   age
##   <int>
## 1    18
## 2    19
## 3    20
## 4    21
## 5    22
## 6    23
## 7    24
## 8    25
## 9    26

```

```
## 10      27
## # ... with 53 more rows
regression_lines <- bind_cols(
  predict(lm_pl_fit, new_data = age_range),
  age_range
)

Wage %>%
  ggplot(aes(age, wage)) +
  geom_point(alpha = 0.2) +
  geom_line(aes(y = .pred), data = regression_lines, color = "blue")
```



All the other variables are categorical. Hot encoding these variables basically results in a “piecewise non-linear” function. However, exploring this most likely is not very interesting. Therefore we will now test splines also on the relationships of age and wage.

```
min(Wage$age)

## [1] 18

spline_recipe <- recipe(formula = wage ~ age, data = Wage_train) %>%
  step_bs(age, options = list(knots = 20, 30, 40, 50, 60, 70))

# Specification
spline_spec <-
  linear_reg() %>%
  set_mode("regression") %>%
  set_engine("lm")
```

```

# Workflow
spline_workflow <- workflow() %>%
  add_recipe(spline_recipe) %>%
  add_model(spline_spec)

# Finalize model
spline_fit <- fit(spline_workflow, data = Wage_train)

# Check RMSE
augment(spline_fit, new_data = Wage_test) %>%
  rmse(truth = wage, estimate = .pred)

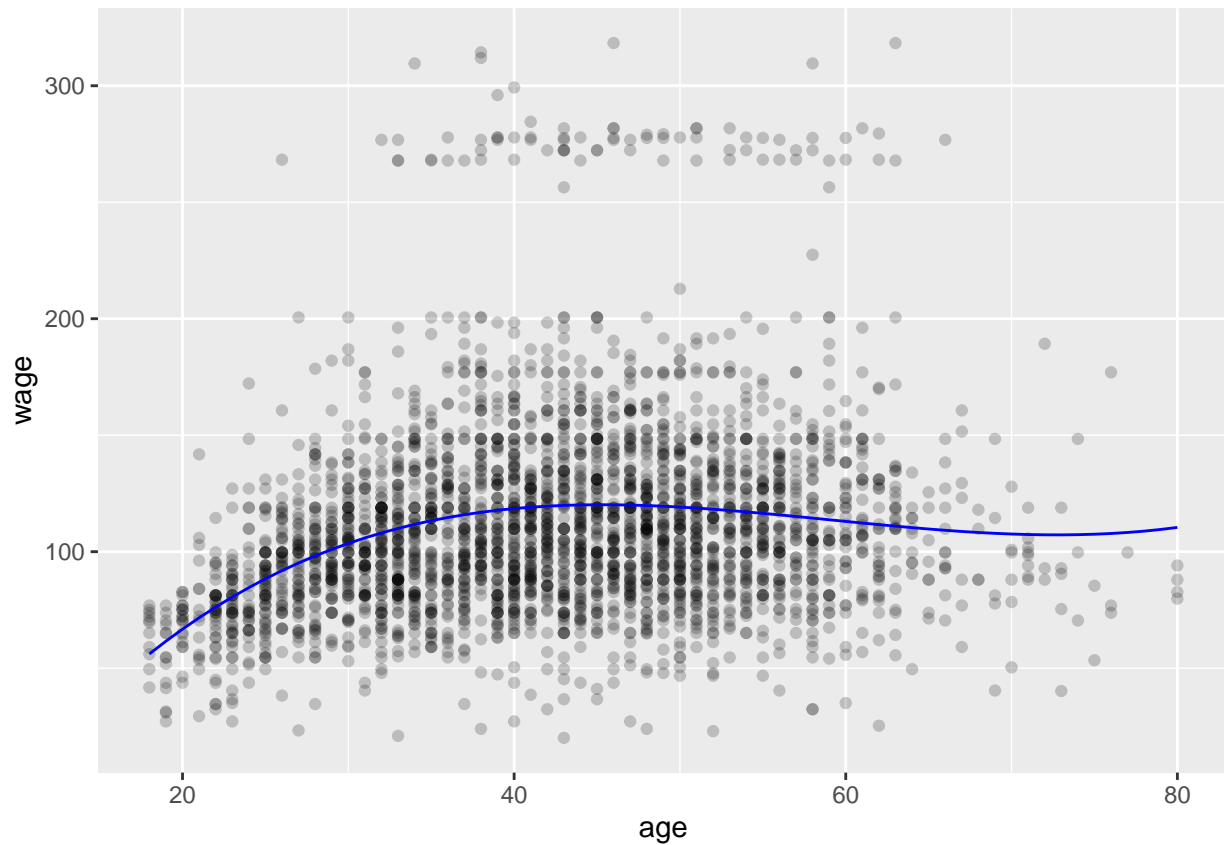
## Warning in bs(x = c(18L, 24L, 43L, 30L, 45L, 34L, 51L, 50L, 56L, 40L, 49L, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard      39.7

regression_lines <- bind_cols(
  predict(spline_fit, new_data = age_range),
  age_range
)

## Warning in bs(x = 18:80, degree = 3L, knots = numeric(0), Boundary.knots =
## c(18L, : some 'x' values beyond boundary knots may cause ill-conditioned bases
Wage %>%
  ggplot(aes(age, wage)) +
  geom_point(alpha = 0.2) +
  geom_line(aes(y = .pred), data = regression_lines, color = "blue")

```

Both the polynomial regression as well as splines capture the relationship between age and wage better than a simple linear regression. We can see that the splines fit the data even better than the polynomial regression, especially at the border regions.